Stoquify moat roadmap

> **30 September 2026 update:** M1-03 local PostgreSQL and EN/FR browser gates passed. Unanswered-request/notification visibility and in-place password confirmation are implemented. Real rail, customer cutover, recipient delivery recovery and pilot gates remain open. See M1-03 execution report; dated 27 September statements below describe the earlier baseline.

Baseline: 27 September 2026. Governing input: moat assessment. Canonical register: MOAT_WORK_PACKAGE_REGISTER.json. This program extends existing owners; it does not certify demand, operations or release readiness.

Keep M1 accountant work desk first; M2 payment-to-close is its first operational use case. M3 is one bounded planning experiment. M4 follows evidence that one supplier relationship is useful. No new customer evidence changes the assessment's priorities. Promote M2 if measured payment pain dominates or firms reject repeat use.

Phase	Planning window	Work and exit gate
P0	Days 1–3	Source baseline, overlap register, dependency contracts
P1	Days 1–14	Prepare onboarding, cutover and outcome measurement; actual reconciled samples are required to pass
P2	Days 15–30+	Accountant triage/detail/acceptance; real selected rail and integrated proof-to-close verification
P3	Days 31–60+	Controlled M1/M2 pilot with actual repeat work, reliability and support costs
P4	Prototype now; fieldwork conditional	Six weekly planning cycles before durable scenario investment
P5	After principal pilot decision	Useful supplier task, repeat participation, then actual cross-buyer overlap

The 90-day horizon is a decision checkpoint, not a release promise. Assume two domain engineers, one frontend engineer, shared QA/design, a product owner, part-time security/SRE and a qualified accountant; staffing is unconfirmed. Complete-scope estimates remain M1 4–8, M2 4–8, conditional M3 8–14 and M4 6–12 engineer-weeks. Provider and expert lead times are additional and unknown. Local changes in this run are smaller than those production scopes.

Critical path: consenting participants and selected rail → reconciled inputs → current authority and recoverable delivery → reviewed resolution and regenerated close proof → repeated productive cycles → sustainable contribution and paid continuation. At most one principal workflow and one bounded experiment are active here.

Existing owners are FA-ACCOUNTANT-ACCESS, FA-ACCOUNTANT-PORTAL, FA-CLOSE-EVIDENCE-ACCEPT, active FA-CASH-COMMAND, PR-17, NTF accountant invitation/delivery and the analytics/network program. This run extends accountant composition; it does not overwrite those registers or payment implementation. Graph navigation uses accountant action community 58, portfolio route community 75 and component community 140 from the August snapshot; current source takes precedence.

Initial discovery: grants, portfolio cards, per-client review and acceptance services, trust exports, reconciliation and supplier proposals exist. The review queue has no action/UI consumer in the inspected tree. Cross-client triage is absent from the existing portfolio. Named rail operations, financial cutover, customer demand and supplier network effects remain unverified. General saved scenario/actual evaluation was not found in the assessed scope.

Execution checkpoint

This run implemented and locally tested M1's submitted-response desk, fixed the client handoff to portal refresh/export, prepared onboarding/productive-work measurement and delivered M3's isolated scenario prototype. M1/M2's full runtime and customer journey is still conditional. Consult the execution report for final verification, rather than interpreting the initial discovery gaps above as the final code state.

Review ownership

Architecture and backend own service composition and domain boundaries; data/migrations own reconciled inputs and rollback; security/privacy own current grants, entitlements and redaction; frontend/workflow design own bilingual states and interaction; QA/release own runtime proof; SRE owns recovery, latency and cost; product/customer success own task baselines, repeat use and paid evidence. Accounting/treasury own reviewed source-to-close correctness. Purchasing/inventory own proposed changes to stock and AP. Country specialists are required if statutory outputs enter the pilot. AI/agent expansion, lending, public ratings and new-country implementation are not applicable to this slice. These are review responsibilities, not claims of independent expert certification.