

Stoquify Breakthrough Landing Experience

Three concepts, scorecard, implementation plan, and verification

RESEARCH REPORT

Prepared 2026-07-27

Evidence-led · bilingual-ready · accessibility-aware · implementation-bounded

Connected operations. Explicit controls. Inspectable evidence.

Research Report and Implementation Plan

Date: 2026-07-27 **Audience:** Stoquify product, design, engineering, growth, and assurance owners **Decision:** Advance a live operating narrative, a role lens, and a claim-to-evidence trust layer—without adding three new landing sections.

1. Executive recommendation

Stoquify already has the right strategic raw material: a premium visual system, a connected operating story, a four-stage early-founder journey, explicit operational boundaries, and unusually strong evidence language. The next advantage will not come from adding more capability cards. It will come from making three things immediately perceptible:

- 1 the system moves as one connected operating model;
- 2 the visitor sees the version of that model relevant to their job;
- 3 every important claim can be inspected through credible evidence.

Implement the following portfolio in order:

- 1 **Foundation — The Proof Ledger:** establish a controlled claim-to-evidence model and upgrade the current trust surface.
- 2 **Lead experience — The Live Operating Twin:** replace a redundant explanatory block with an interactive but server-readable operating-event narrative.
- 3 **Supporting experiment — The Role-Lens Command Brief:** let visitors select a role and reframe the same canonical product truth without multiplying pages or promises.

The **Live Operating Twin** is the visual and messaging breakthrough. It turns Stoquify's differentiator—connected, controlled operations—into a behavior the visitor can see. The **Role Lens** creates personal relevance and a continuity signal for onboarding. The **Proof Ledger** makes the experience professionally credible and protects the team from marketing claims that outrun implementation evidence.

This portfolio should consolidate the current page rather than increase its 12-section length. It should retain the qualified `/#pricing` rollout CTA until adoption and reserve `/register` for the adoption decision.

2. Evidence base and method

The assessment combined:

- 1 inspection of Stoquify's landing route and its 12 current top-level components;
- 1 review of the hero, connected workflow, product gallery, founder journey, use cases, trust, adoption/pricing, and final CTA;
- 1 review of English and French message keys;
- 1 review of recent landing strategy and UI/UX reports;
- 1 inspection of focused structural and browser certification for the founder journey;
- 1 architecture-graph context for the connected operations story, treated as directional because its snapshot predates current landing work;
- 1 study of Accio's official localized founder-oriented experience;
- 1 comparison with eleven additional official product sites.

Current landing sequence

The route currently renders:

- 1 Hero
- 2 Disconnect Problem
- 3 Product Gallery
- 4 Connected Workflow
- 5 Operations Map
- 6 People to Pay
- 7 Module Deep Dives
- 8 Automation
- 9 Trust
- 10 Use Cases
- 11 Adoption/Pricing
- 12 Final CTA

Recent Stoquify reviews correctly identify the opportunity and constraint: the page is differentiated and visually strong, but the proof is distributed across a dense sequence. New work should subtract, merge, or replace before it adds.

Completed work that must not be re-proposed

The four-stage early-founder journey is already present inside the Operations Map:

- 1 foundation;
- 2 launch;
- 3 control;
- 4 growth.

It is bilingual, ordered rather than carousel-based, and connected to the qualified `/#pricing` rollout CTA. Focused tests verify English/French stage equivalence, sequence, destination, and bounded claims. Browser evidence covers English and French at mobile, tablet, and desktop widths, with reduced motion, keyboard focus, overflow, clipping, overlap, page-error, console-error, and request-failure checks passing.

Therefore, “add a founder journey” is not a new idea. This report treats that journey as the narrative baseline and proposes the next presentation layer.

3. What Accio’s early-stage-founder experience gets right

Accio’s official localized surfaces present an early-stage founder with a recognizable progression from idea and market research through sourcing, product development, a website, and first-sale readiness. The useful pattern is not the exact feature set; it is the way the page converts a broad platform into a guided sequence.

Transferable mechanics

- 1 **Self-identification before explanation.** The visitor can recognize “early-stage founder” as their context rather than parse an undifferentiated feature inventory.
- 1 **Progress as the organizing metaphor.** Capabilities appear as steps toward an outcome, reducing cognitive load.
- 1 **A team-like system.** The product is framed as coordinated assistance, not isolated tools.

- 1 **Visible momentum.** The interaction suggests movement from an incomplete state to a completed operating outcome.
- 1 **Stage-relevant action.** Calls to action fit what the founder is trying to accomplish now.

What Stoquify should not copy

- 1 Stoquify should not imply autonomous execution where approvals, evidence, or statutory responsibility remain human-controlled.
- 1 It should not translate an idea-to-product workflow literally; Stoquify’s advantage begins when operational events must stay connected, scoped, reviewable, and financially trustworthy.
- 1 It should not rely on novelty animation to carry meaning.
- 1 It should not turn every module into a simulated “agent” without a stable responsibility boundary.

The Stoquify adaptation

Accio helps a founder see how an idea becomes a business. Stoquify should help an operator see how a business event becomes **controlled operational truth**:

event → control → evidence → financial position → decision

That chain is more defensible, more distinctive, and closer to the product’s real architecture.

4. Cross-site pattern library

The comparison set was selected for product-led storytelling, complex-system explanation, role relevance, trust, or interaction quality. The goal was to extract mechanics, not imitate layouts.

Accio	Persona-led, staged path from idea to operating outcome	Preserve the completed founder journey; carry the visitor into a live operating event	Broad autonomous-agent implications
Stripe	One durable growth promise, flexible paths, proof at scale	State one operating promise; keep depth behind progressive disclosure	Scale claims without Stoquify evidence
Linear	Numbered workflow narrative with product UI doing the explaining	Make the event chain the visual spine	Aesthetic minimalism without accounting context
Ramp	Outcome-plus-control storytelling for finance operations	Place prevention, evidence, and decision next to each event	Unverified savings claims
Mercury	Clear founder language, product grouping, and demo cues	Use founder-recognizable jobs and a bounded “see it operate” cue	Banking-category language
Vercel	Live system behavior and audience-oriented product framing	Show state transition rather than static screenshots alone	Performance-heavy spectacle
Notion	Progressive use-case previews and a platform that adapts to work	Reframe canonical truth by role	Endless template breadth
Shopify	A simple lifecycle from starting to selling and growing	Keep the founder journey outcome-led	“Start instantly” language where configuration is required
Brex	Role/segment paths combined with control and quantified proof	Pair role relevance with inspectable claims	Customer outcomes not owned by Stoquify
Deel	Visitor chooses a job; a today-to-tomorrow progression makes speed legible	Offer a role/job lens and staged operating brief	Compliance universality

Attio	Product behavior is embedded directly into the narrative; context improves relevance	Demonstrate the system acting on connected facts	"Always-on" autonomy beyond guarded capabilities
Retool	Interactive construction plus governance positioned together	Present power and control as a single experience	Builder metaphors that obscure the packaged product

Combined lesson

The best current pages combine three layers:

- 1 a compact category promise;
- 2 a product behavior the visitor can see;
- 3 proof that reduces the risk of believing the promise.

Stoquify currently has all three ingredients, but they are spread across many sections. The proposed concepts concentrate them.

5. Concept One — The Live Operating Twin

One-line idea

Let a visitor choose one familiar business event and watch Stoquify carry it through control, evidence, accounting context, and the next decision.

Founder problem

An early-stage founder can understand "POS," "inventory," "payroll," and "accounting" individually but may not understand why a connected operating system matters. Static module cards reinforce the impression of a suite of tools. The product's actual moat—the integrity of the transitions between tools—remains abstract.

Core experience

The page offers four bounded sample events:

- 1 a sale is completed;
- 1 stock is received;
- 1 a supplier payment is reviewed;
- 1 payroll inputs reach readiness.

Selecting an event updates a single operating canvas:

- 1 **Source event** — who, where, when, and what happened;
- 2 **Control** — validation, permission, variance, or approval state;
- 3 **Evidence** — linked document, movement, provider, or reviewer trail;
- 4 **Financial effect** — the relevant ledger or obligation context, clearly labelled as representative where necessary;
- 5 **Next decision** — the action an owner, operator, or accountant must take.

The visual should feel like a calm control-room trace, not a dashboard montage. Lines or state transitions may animate only after direct interaction and must collapse to an ordered, fully understandable sequence under reduced motion.

Why it is distinctively Stoquify

Many products can display a polished dashboard. Fewer can demonstrate that an operational event retains tenant, branch, actor, source, control state, and evidence as it influences financial understanding. This concept makes Stoquify's evidence-first operating architecture visible without requiring the visitor to know accounting vocabulary.

Inspired by

- 1 Accio's guided progress;
- 1 Linear's product-led numbered flow;
- 1 Vercel's visible system behavior;
- 1 Ramp's control-and-outcome pairing;
- 1 Attio's narrative built from product activity.

Desktop treatment

- 1 A left-side event selector with four clearly labelled buttons.
- 1 A central five-step evidence chain.
- 1 A right-side "owner brief" showing what changed, what is blocked, and what needs attention.
- 1 One qualified CTA: "Plan this workflow" linking to `/#pricing`.
- 1 No autoplay. The first event is server-rendered as the default.

Mobile treatment

- 1 Event buttons become a two-by-two grid or compact select-plus-tabs.
- 1 The chain becomes a vertical ordered list.
- 1 The owner brief follows the chain rather than floating.
- 1 Sticky behavior is avoided unless testing proves it does not obscure content.

Accessibility and motion

- 1 Use real buttons with `aria-pressed` or a tab pattern only if the content semantics fit tabs.
- 1 Keep every stage in DOM order.
- 1 Move focus only when the visitor explicitly requests it.
- 1 Provide a concise status announcement when an event changes.
- 1 Under `prefers-reduced-motion`, switch states without animated travel.
- 1 Do not encode control state by color alone.

Minimum change to Stoquify

Replace or consolidate the current **Connected Workflow** and part of **Product Gallery** into one live narrative. Keep the existing Founder Journey in Operations Map as the lifecycle overview. The journey answers "where am I going?"; the Operating Twin answers "how does Stoquify work?"

Recommended surface ownership:

- 1 server component: default narrative, labels, boundaries, and CTA;
- 1 small client island: event selection and progressive state transition;
- 1 localized content: `landing.operatingTwin`;
- 1 typed sample-event registry: presentation-only facts and evidence labels;

- 1 no tenant API calls on the public page.

Content and data dependencies

- 1 four approved representative event stories;
- 1 verified links between each story and implemented product surfaces;
- 1 clearly labelled sample amounts and identities;
- 1 screenshots or UI fragments approved for public display;
- 1 English and French parity.

Boundaries

- 1 No live tenant data.
- 1 No autonomous action claim.
- 1 No claim that a financial effect is posted unless the demonstrated product path really posts it.
- 1 No implication of statutory certification.
- 1 Samples must be labelled “representative” or “sample.”

Expected impact

Attention and comprehension

- 1 A meaningful interaction gives the eye a job and makes the connected-system story memorable.
- 1 The event chain teaches the product model faster than another catalogue.

Qualified conversion

- 1 A visitor can select the workflow closest to their risk and reach adoption with context.
- 1 Sales or onboarding can receive only a non-sensitive workflow intent after explicit submission.

Activation and retention potential

- 1 The selected workflow can become the suggested first rollout path after consent.
- 1 The same five-step mental model can reappear in onboarding and the authenticated command center, improving continuity.

Measurement

Leading indicators:

- 1 operating-event selection rate;
- 1 completion rate across the five-step chain;
- 1 “Plan this workflow” click-through by event;
- 1 task-test accuracy for “what makes Stoquify different?”

Lagging indicators:

- 1 qualified rollout starts by selected workflow;
- 1 completion of the first corresponding onboarding milestone;
- 1 week-two return to the corresponding operational surface.

Falsification:

Stop or simplify if visitors interact but cannot explain the event-to-evidence relationship better than the static baseline, or if the client enhancement harms Core Web Vitals, keyboard flow, or mobile comprehension.

Delivery estimate

Medium. The major work is content verification and interaction design, not backend integration.

6. Concept Two — The Role-Lens Command Brief

One-line idea

Allow a visitor to view one canonical operating system through the daily decisions of a founder, branch manager, accountant, or people/operations lead.

Founder problem

Stoquify spans several responsibilities. A founder wants control without operating every task; an accountant wants source-linked numbers; a branch manager wants exceptions and ownership; a people lead wants approved inputs and clear boundaries. A single dense narrative can feel comprehensive but impersonal.

Core experience

A compact role selector appears near the beginning of the product story:

- 1 Founder / Owner
- 1 Branch Manager
- 1 Accountant
- 1 People & Operations

The selection does not change product truth. It changes:

- 1 the lead question;
- 1 the order of three relevant decisions;
- 1 the highlighted evidence;
- 1 the most relevant sample event;
- 1 the label of the qualified rollout CTA.

Example founder brief:

- 1 What needs my decision today?
- 1 Which branch or workflow is drifting?
- 1 Can I trust the number behind this signal?

Example accountant brief:

- 1 Which source produced this balance?
- 1 What remains unreconciled or unreviewed?
- 1 What evidence will travel with the close?

The selected role persists in the URL query or a short-lived, non-sensitive local preference. It should only enter an account or CRM after the user explicitly submits it.

Why it is distinctively Stoquify

Most role personalization changes benefit copy. Stoquify can change the **decision-and-evidence lens** while preserving one controlled source of truth. That demonstrates a product principle: different roles see appropriate work, but the underlying facts do not fork.

Inspired by

- 1 Accio's persona entry;
- 1 Deel's job-first chooser;
- 1 Mercury's founder-focused pathways;
- 1 Notion's adaptable use-case previews;
- 1 Brex's role and segment routing.

Desktop treatment

- 1 A small, visible role rail under the category promise or at the start of the consolidated product story.
- 1 The selected role changes one command-brief card and emphasis inside the Operating Twin.
- 1 A concise "Same operating truth, relevant decisions" explanation prevents the selector from looking like separate products.

Mobile treatment

- 1 Four wrap-safe buttons or an accessible native/select-enhanced control.
- 1 Role summary appears immediately below.
- 1 No sideways role carousel.
- 1 Every role remains reachable with keyboard and touch without precision gestures.

Accessibility and motion

- 1 Use buttons with a persistent selected state.
- 1 Announce the new brief without stealing focus.
- 1 Ensure meaningful content remains available when JavaScript is unavailable: Founder / Owner is the server-rendered default, with links to role-specific query URLs as an enhancement option.
- 1 Keep text length comparable across English and French to minimize layout shifts.

Minimum change to Stoquify

Use the Role Lens to consolidate the current **Use Cases** and buyer-path portions of **Adoption/Pricing**, not to create another section. The existing pricing/adoption model remains authoritative.

Recommended surface ownership:

- 1 server-rendered role content registry;
- 1 lightweight client selection;
- 1 URL-search-param support for shareable role views;
- 1 optional local preference with no identity or tenant information;
- 1 no entitlement, permission, or product-surface behavior change.

Content and data dependencies

- 1 product-owner-approved jobs and decision language;

- 1 mapping from each role to existing, verifiable product capabilities;
- 1 consistent English and French strings;
- 1 analytics taxonomy that records only role choice and content interaction.

Boundaries

- 1 A role selection is marketing context, not authorization.
- 1 Do not infer employment, company, geography, or financial state.
- 1 Do not preselect a role from opaque tracking.
- 1 Do not personalize statutory claims by jurisdiction without an approved country-pack source.
- 1 Never let the selector imply access to surfaces the product would deny after login.

Expected impact

Message relevance and attention

- 1 Visitors recognize their daily problem sooner.
- 1 The same broad system becomes easier to scan without shrinking its real scope.

Qualified conversion

- 1 CTA language can carry the visitor's selected operating concern into an explicit rollout request.
- 1 Sales conversations begin with a job and workflow rather than an undifferentiated "demo."

Activation and retention potential

- 1 With consent, the chosen lens can prioritize onboarding guidance.
- 1 Reusing the same lens in authenticated daily briefs supports habit formation around role-appropriate decisions.

Measurement

Leading indicators:

- 1 role-selection rate;
- 1 role-brief completion or expansion rate;
- 1 qualified CTA rate by role;
- 1 first-click accuracy in role-based comprehension tests.

Lagging indicators:

- 1 onboarding milestone completion by submitted role;
- 1 reduction in time to first relevant workflow;
- 1 week-two return to role-relevant command surfaces.

Falsification:

Remove or reduce the selector if most visitors ignore it, if role-specific copy does not improve task comprehension or qualified CTA behavior, or if the mapping creates product expectations that cannot be honored.

Delivery estimate

Medium. Content governance and analytics interpretation are the main risks.

7. Concept Three — The Proof Ledger

One-line idea

Turn major landing claims into inspectable, dated claim-to-evidence records instead of relying on generic trust language.

Founder problem

Operational and financial software asks for a high level of trust. Early-stage companies may not yet have a large public customer-proof library, and Stoquify must not fill that gap with vague compliance language or invented outcomes. Its existing landing copy contains strong evidence concepts, but the visitor cannot consistently distinguish:

- 1 implemented control;
- 1 representative sample;
- 1 supported workflow;
- 1 planned capability;
- 1 assurance boundary.

Core experience

Each high-value claim receives a compact proof marker. Opening it reveals:

- 1 **Claim:** the exact statement being supported;
- 1 **Status:** available, controlled pilot, representative sample, or planned—not a vague badge;
- 1 **Product evidence:** screenshot, route, test, or workflow reference approved for public use;
- 1 **Scope:** module, actor, location, and jurisdiction limits where relevant;
- 1 **Control:** permission, approval, reconciliation, audit, or provenance behavior;
- 1 **Last reviewed:** date and accountable content owner;
- 1 **Boundary:** what the claim does not mean.

The Trust section becomes a curated public proof index rather than a long list of assurances. A future `/trust` route may provide the full register, but the landing page exposes only the few records needed to support its central promise.

Why it is distinctively Stoquify

Stoquify's product architecture repeatedly emphasizes source, scope, actor, evidence state, and review history. Applying the same discipline to marketing creates unusual coherence: the company presents its product claims the way the product presents operational truth.

Inspired by

- 1 Stripe's proof adjacent to platform claims;
- 1 Ramp and Brex's control-and-outcome evidence;
- 1 Retool's governance integrated with product power;
- 1 Attio and Notion's customer and update proof;
- 1 Stoquify's existing assurance and evidence language.

Desktop treatment

- 1 A small proof glyph beside no more than five primary claims.
- 1 A side panel or inline disclosure with a consistent evidence schema.

- 1 A Trust-section index grouped by “Operational trace,” “Access and approval,” “Financial evidence,” and “Product status.”
- 1 No badge wall without explanatory text.

Mobile treatment

- 1 Inline accordion disclosures.
- 1 Short claim title first, evidence details second.
- 1 Large close target and reliable back navigation.
- 1 Avoid modal traps for long evidence.

Accessibility and motion

- 1 Native disclosure patterns or correctly implemented dialog semantics.
- 1 Clear status text in addition to icons and color.
- 1 Evidence images require purposeful alternative text.
- 1 Links describe destination and evidence type.
- 1 No animated counters or auto-rotating testimonials.

Minimum change to Stoquify

Upgrade the existing **Trust** component and connect proof markers to the consolidated Operating Twin. Do not create a new landing section. Introduce a typed claim registry that separates:

- 1 public wording;
- 1 evidence type and source;
- 1 status;
- 1 scope and boundary;
- 1 locale strings;
- 1 last-review metadata.

A public `/trust` route is a separate, later decision. The minimum landing version can remain entirely local and server-rendered.

Content and data dependencies

- 1 an approved inventory of five or fewer primary claims;
- 1 evidence owner for each claim;
- 1 public-safe screenshots, tests, or route references;
- 1 review dates and status vocabulary;
- 1 a process for withdrawing stale proof.

Boundaries

- 1 No self-certification.
- 1 No legal or statutory assurance beyond owned evidence.
- 1 No confidential test, tenant, security, or architecture data.
- 1 A passing internal test is evidence of behavior, not market impact.
- 1 Planned capability is visibly distinct from implemented capability.

- 1
- Customer outcomes require customer approval and a traceable source.

Expected impact

Professional trust and differentiation

- 1
- The presentation demonstrates maturity through explicit evidence and limitations.
- 1
- Stoquify can build credibility before it has a large testimonial library.

Conversion quality

- 1
- Serious buyers can inspect the control behind the message before requesting a rollout.
- 1
- Product and sales teams share a safer source for public claims.

Retention and expansion potential

- 1
- The proof vocabulary can continue into onboarding, product evidence views, release notes, and module expansion decisions.
- 1
- Consistent claim status reduces expectation gaps that create early churn.

Measurement

Leading indicators:

- 1
- proof-marker open rate by claim;
- 1
- completion of a proof record;
- 1
- trust-to-rollout CTA sequence;
- 1
- usability-test accuracy distinguishing implemented, sample, and planned states.

Lagging indicators:

- 1
- change in assurance or scope objections during qualified conversations;
- 1
- reduction in expectation-related onboarding escalations;
- 1
- module-expansion engagement with relevant proof.

Falsification:

Simplify or relocate the Proof Ledger if visitors cannot understand the status vocabulary, if disclosures distract from the primary narrative, or if evidence owners cannot keep records current.

Delivery estimate

Low-to-medium for the landing registry; higher only if a public Trust Center and automated evidence publication are added later.

8. Comparative scorecard

Scale: 1 = weak/unfavorable; 5 = strong/favorable. For delivery feasibility and performance safety, 5 means easier or safer.

Differentiation	5	4	5
Message clarity	5	5	4
Visual memorability	5	4	4

Professional trust	4	4	5
Qualified conversion potential	4	5	4
Activation and retention potential	4	5	4
Mobile resilience	4	5	5
Accessibility	4	4	5
Delivery feasibility	3	3	4
Performance safety	4	5	5
Current-architecture fit	5	4	5
Total / 55	47	48	50

The totals do not determine sequencing by themselves. The Proof Ledger scores highest because it is safe and foundational. The Operating Twin is the lead creative bet because it provides the most visible breakthrough. The Role Lens has the strongest activation hypothesis but depends on a stable canonical narrative and careful measurement.

9. Recommended portfolio and page architecture

Recommended hierarchy

- 1 **Category promise:** connected operations become controlled, reviewable truth.
- 2 **Founder lifecycle:** the existing foundation-to-growth journey.
- 3 **Live behavior:** one event moving through the Operating Twin.
- 4 **Role relevance:** the same truth reframed as a role-specific command brief.
- 5 **Inspectable proof:** proof records adjacent to central claims.
- 6 **Adoption decision:** choose the rollout path, then register.

Consolidation target

Do not append three sections. Instead:

- 1 retain Hero;
- 1 retain and tighten the Disconnect Problem;
- 1 merge Product Gallery and Connected Workflow into the Live Operating Twin;
- 1 retain Operations Map and its completed Founder Journey;
- 1 preserve only the most valuable deep-dive content;
- 1 use Role Lens to consolidate Use Cases and buyer-path repetition;
- 1 upgrade Trust into the Proof Ledger;
- 1 retain Adoption/Pricing as the explicit commercial decision;
- 1 retain Final CTA only if it adds a distinct, non-repetitive action.

This moves toward the existing nine-section aspiration while increasing perceived product depth.

10. Phased implementation plan

Phase 0 — Freeze the baseline and claims

Goal: agree on the current truth before designing the new interaction.

Actions:

- 1 snapshot the current landing structure, copy, test suite, and browser evidence;
- 1 identify the five primary public claims the page must prove;
- 1 classify each claim as implemented, controlled pilot, representative sample, or planned;
- 1 name a product/evidence owner and content owner for each;
- 1 approve the four representative operating events;
- 1 confirm the four role lenses and their exact jobs.

Acceptance criteria:

- 1 no claim lacks an accountable owner, evidence source, scope, and boundary;
- 1 the existing founder-journey tests remain green;
- 1 English and French content responsibilities are assigned;
- 1 no landing concept requires new tenant entitlements or broad authorization enforcement.

Phase 1 — Build the Proof Ledger foundation

Goal: make public claims safe and reusable before increasing interaction.

Actions:

- 1 define a typed claim registry;
- 1 implement proof markers for two pilot claims;
- 1 upgrade the Trust component into a compact proof index;
- 1 add status, scope, last-reviewed, and boundary copy;
- 1 add structural tests preventing missing or stale-required fields;
- 1 complete product, legal/compliance, and security review of exposed evidence.

Acceptance criteria:

- 1 primary text renders on the server;
- 1 keyboard and screen-reader users can open and close proof records;
- 1 status is not color-only;
- 1 no private route, tenant identifier, customer data, or confidential control detail is exposed;
- 1 every statement has English/French parity or an explicit release block.

Phase 2 — Prototype and ship the Live Operating Twin

Goal: make the connected operating model visible.

Actions:

- 1 choose one highest-confidence event for the server-rendered default;
- 1 design the five-step chain and owner brief;
- 1 create a lightweight client selector for four events;

- 1 replace or merge redundant Product Gallery / Connected Workflow content;
- 1 attach proof records to the two most important event claims;
- 1 route the CTA to `/#pricing` with bounded workflow context;
- 1 instrument anonymous, non-sensitive event selection and CTA behavior.

Acceptance criteria:

- 1 meaningful content and CTA remain usable without JavaScript;
- 1 no autoplay, scroll-jacking, or forced carousel;
- 1 reduced-motion mode removes animated travel;
- 1 no mobile overflow or clipped evidence;
- 1 no critical page, console, or request errors;
- 1 the interaction bundle meets the agreed performance budget;
- 1 comprehension testing beats the static baseline before full rollout.

Phase 3 — Run the Role-Lens experiment

Goal: test whether job relevance improves comprehension and qualified action.

Actions:

- 1 implement Founder / Owner as the default;
- 1 add Branch Manager, Accountant, and People & Operations views;
- 1 change ordering and emphasis, not underlying capability truth;
- 1 support shareable query-state URLs;
- 1 optionally retain a non-sensitive local role preference;
- 1 carry the role into onboarding only after explicit submission or authenticated confirmation;
- 1 test the personalized version against the canonical baseline.

Acceptance criteria:

- 1 role choice is never treated as authorization;
- 1 each role maps only to implemented or clearly bounded capability;
- 1 no opaque inference or personal profiling;
- 1 role selection is keyboard operable and announced correctly;
- 1 the experiment has predeclared success and stop thresholds.

Phase 4 — Connect acquisition to activation

Goal: determine whether the landing mental model survives first use.

Actions:

- 1 reuse the chosen operating-event language in rollout planning;
- 1 reuse role-appropriate brief language in onboarding;
- 1 measure first relevant workflow completion;
- 1 compare week-two return and expectation-related support events;
- 1 retire variants that attract clicks but reduce qualified activation;

- 1 update proof records on a fixed review cadence.

Acceptance criteria:

- 1 analytics contain no operational payload or sensitive tenant facts;
- 1 conversion is not declared a success without downstream activation quality;
- 1 stale proof is withdrawn or visibly reclassified;
- 1 the product owner signs off before any claim graduates from pilot/sample to available.

11. Test and certification plan

Structural tests

- 1 exactly four operating events;
- 1 exactly five semantic stages per event;
- 1 one qualified pre-adoption CTA destination;
- 1 no direct registration CTA before adoption;
- 1 English/French key parity;
- 1 no unsupported promise patterns;
- 1 every published claim has status, source type, scope, boundary, owner, and review date;
- 1 role views reference the same canonical capability identifiers.

Component tests

- 1 event and role selection;
- 1 selected and focus state;
- 1 status announcements;
- 1 disclosure behavior;
- 1 query-state hydration;
- 1 no-JavaScript default;
- 1 stale or missing proof handled safely.

Browser matrix

Test English and French at:

- 1 390 × 844 mobile;
- 1 768 × 1024 tablet;
- 1 1440 × 1000 desktop.

For each:

- 1 default and one alternate event;
- 1 default and one alternate role;
- 1 standard and reduced motion;
- 1 keyboard-only path;
- 1 proof disclosure;

- 1 CTA destination;
- 1 horizontal overflow;
- 1 clipping and overlap;
- 1 page and console errors;
- 1 critical request failures.

Human comprehension tests

Ask a founder, operator, and accountant to answer:

- 1 What does Stoquify connect?
- 2 What happens after a business event?
- 3 What must a person still approve or review?
- 4 Which evidence supports the claim?
- 5 What is the next action if you want to adopt it?

The concept should not ship broadly if interaction increases but correct answers do not.

12. Measurement framework

Attention

- 1 meaningful event or role interaction;
- 1 completion of the operating trace;
- 1 scroll progression into adoption;
- 1 return to the product narrative after opening proof.

Comprehension

- 1 task-test correctness;
- 1 time to state the connected-system differentiator;
- 1 ability to distinguish sample, pilot, and available;
- 1 ability to identify the human decision point.

Qualified conversion

- 1 rollout CTA by selected event and role;
- 1 adoption-path completion;
- 1 qualified submission rate;
- 1 rate of mismatched or unsupported expectations.

Activation

- 1 first relevant workflow configured;
- 1 first evidence-bearing event completed;
- 1 time to first role-relevant control or decision;
- 1 onboarding completion quality, not only account creation.

Retention

- 1 week-two return to the relevant workflow;
- 1 repeat use of evidence or command-brief surfaces;
- 1 expectation-related support incidents;
- 1 expansion into a second connected workflow.

Experiment discipline

- 1 pre-register one primary metric and two guardrails per experiment;
- 1 compare against the current certified baseline;
- 1 segment only by explicit, non-sensitive selections;
- 1 do not optimize for dwell time if comprehension or activation declines;
- 1 stop variants that harm accessibility, performance, or trust.

13. Risks and guardrails

Risk: the page becomes even denser

Mitigation: require every new concept to replace or consolidate an existing block. The Operating Twin absorbs static workflow repetition; Role Lens absorbs use-case repetition; Proof Ledger upgrades Trust.

Risk: polished samples appear to be live product data

Mitigation: label representative data at the visual and record level. Avoid realistic customer identities. Never use tenant data.

Risk: personalization becomes profiling

Mitigation: use explicit selection, store no sensitive attribute, and require consent before carrying context into an account or CRM.

Risk: evidence language becomes self-certification

Mitigation: distinguish behavior evidence from regulatory assurance. Publish scope and boundary next to status. Require accountable review.

Risk: client interaction weakens performance or accessibility

Mitigation: server-render the default, isolate a small enhancement, avoid autoplay, certify reduced motion and keyboard use, and enforce a bundle budget.

Risk: marketing truth diverges from product truth

Mitigation: use typed canonical capability and claim identifiers, review dates, and withdrawal rules. Role views may reorder truth but cannot invent it.

Risk: retention is overclaimed

Mitigation: present retention as a hypothesis. Measure whether the selected mental model improves activation and repeat use; do not infer causality from landing engagement alone.

14. Decisions requiring owner approval

Before implementation, product owners must approve:

- 1 the central category promise;
- 1 the four public sample events;
- 1 the four public role lenses;
- 1 the five primary claims;
- 1 the status vocabulary;
- 1 which screenshots or tests are safe to expose;
- 1 the performance budget;
- 1 the measurement consent and retention policy;
- 1 whether a public /trust route belongs in this release or a later one.

No broad entitlement-enforcement change is required or recommended by this plan.

15. Source register

Stoquify sources reviewed

- 1 current landing route composition and landing components;
- 1 English and French landing messages;
- 1 focused founder-journey structural tests;
- 1 founder-journey browser evidence for English/French and mobile/tablet/desktop;
- 1 STOQUIFY_LANDING_PAGE_COMPETITIVE_REVIEW_2026-07-19.md;
- 1 AQSTOQFLOW_UIUX_PHASE_07_PUBLIC_FIRST_IMPRESSION_REPORT_2026-07-18.md;
- 1 landing-packages-adoption-report-2026-07-18.md;
- 1 STOQUIFY_LANDING_PAGE_FINAL_9_PLUS_COMPLETION_PLAN_2026-07-20.md;
- 1 landing-related architecture graph/report context.

Official external sources

- 1 [Accio](#)
- 1 [Accio Work — Funding](#)
- 1 [Stripe](#)
- 1 [Linear](#)
- 1 [Linear for Startups](#)
- 1 [Ramp Platform](#)
- 1 [Mercury](#)
- 1 [Vercel](#)
- 1 [Notion Product](#)
- 1 [Shopify — Start](#)
- 1 [Brex](#)

1 [Deel](#)

1 [Attio](#)

1 [Retool](#)

External sites change over time. These sources were reviewed on 2026-07-27; patterns are observations, not endorsements or benchmarks.

16. Final recommendation

Stoquify should not try to win by looking like a larger generic suite. It should win by making its operating philosophy visible:

A business event remains connected to its scope, control, evidence, financial meaning, and next decision.

The **Live Operating Twin** makes that proposition captivating. The **Role-Lens Command Brief** makes it personally relevant. The **Proof Ledger** makes it credible and professionally defensible.

Build the proof foundation first, use it to support one high-quality Operating Twin, and then test role relevance as a bounded enhancement. That sequence gives Stoquify a stronger presentation, clearer messaging, better attention mechanics, and a plausible bridge from acquisition to retention—without adding feature sprawl or claiming more than the product can prove.