STOQUIFY

Top 12 Additive Agents and Skills — Research Prompt

Execution-ready brief for a repository-grounded and internet-grounded Stoquify capability review

This document preserves the complete prompt used to evaluate professional agents, reusable skills, and Copilot capabilities for daily use, retention, referrals, commercial value, and defensible platform moats.

Prepared 2 August 2026

Repository-grounded • Primary-source researched • Human-controlled

Act as a multidisciplinary principal review team: system architect, cyber-security architect, senior frontend engineer, UI/UX specialist, product strategist, business logic analyst, and SaaS growth advisor.

Additionally, act as:

- A principal AI-agent and multi-agent systems architect.
- An enterprise workflow-automation specialist.
- An AI product and platform strategist.
- A SaaS retention, product-led growth, and referral-loop specialist.
- An African fintech, accounting, compliance, and SMB operations specialist.
- A developer-experience and agent-skills ecosystem architect.
- An AI security, governance, evaluation, and observability specialist.

Mission

Conduct a rigorous, evidence-led investigation to identify the 12 most valuable professional AI agents, agent archetypes, reusable skills, or specialized copilot capabilities that Stoquify should consider adding to its Copilot ecosystem.

The recommendations must help Stoquify become an indispensable, trustworthy daily operating platform for its target users—not merely add attractive AI features.

Each recommended agent or skill must materially advance at least one of these outcomes:

1. Increase daily and hourly active use.
2. Increase user retention and workflow dependency.
3. Create natural referral, invitation, or collaboration loops.
4. Reduce the time, skill, or effort required to operate a business.
5. Improve financial control, operational visibility, compliance, and decision-making.
6. Strengthen Stoquify’s proprietary data, workflow, evidence, integration, or intelligence moats.
7. Advance Stoquify’s primary goal of becoming an evidence-backed business operating and financial control system for African businesses and their professional partners.

Critical Interpretation

“Agents and skills” may include:

- Specialized autonomous or assisted agent roles.
- Reusable professional skills and workflow packages.
- Agent orchestration and governance capabilities.
- Domain-specific copilots.
- Monitoring and intelligence agents.
- Collaboration, reporting, compliance, and advisory agents.
- Agent-supporting infrastructure that directly unlocks valuable user-facing capabilities.

Do not recommend generic chatbots or capabilities simply because they are fashionable.

Do not assume that an externally available agent package should be installed directly. Treat external implementations as research evidence and inspiration until their licensing, security, maintenance, compatibility, and quality have been assessed.

Phase 1 — Understand Stoquify First

Before searching the internet, inspect the Stoquify repository and its existing documentation to establish:

- The platform's expressed mission and target users.
- Existing product modules, workflows, and service boundaries.
- The current Copilot backbone and proposed runtime architecture.
- Existing agents, skills, evaluation cases, registries, and shared contracts.
- Already proposed or installed capabilities.
- Architectural, security, data, compliance, and execution constraints.
- Major gaps that prevent Stoquify from becoming a daily operating system.
- Opportunities for recurring workflows and multi-user collaboration.

Inspect relevant evidence under:

- `docs/copilot/`
- `docs/agents and skills/`
- `graphify-out/`
- `app/`
- `services/`
- `actions/`
- `components/`
- `hooks/`
- `lib/`
- `config/`
- `prisma/`
- Existing agent and skill registries
- Existing validation and evaluation suites

Use the architectural knowledge graphs in `graphify-out/` when assessing system relationships, integration points, duplication, and implementation impact.

Explicitly identify the existing agent and skill inventory so the final recommendations do not duplicate capabilities that Stoquify already possesses.

Phase 2 — Conduct Current Internet Research

Search the current internet for professional agents, agent frameworks, skills, plugins, enterprise copilots, workflow products, open-source projects, research, and proven SaaS patterns relevant to Stoquify.

Prioritize:

- Official product documentation.
- Maintainer-owned repositories.
- Framework documentation and release notes.
- Reputable technical research.
- Demonstrated enterprise implementations.
- Credible product benchmarks and case studies.
- Actively maintained projects with clear licensing.
- Solutions relevant to African business operations, OHADA markets, emerging-market infrastructure, and multilingual or low-connectivity environments.

Research areas should include, where relevant:

- Accounting and financial operations.
- Cash-flow management and forecasting.
- Payment and bank reconciliation.
- Inventory and purchasing.
- Payroll and workforce operations.
- Compliance and close assurance.
- Business intelligence and decision support.
- Customer and supplier collaboration.
- Anomaly, fraud, and risk detection.
- Document and evidence intelligence.
- Workflow orchestration and exception management.
- Professional-advisor and accountant portals.
- Embedded support and user education.
- Data integration and connector intelligence.
- Agent evaluation, governance, and observability.

For every important factual claim, provide a direct source link and access date. Prefer primary sources. Distinguish verified facts from your own architectural or commercial inferences.

Reject or penalize candidates that are:

- Abandoned or poorly maintained.
- Primarily marketing demonstrations.
- Insecure or excessively permissioned.
- Incompatible with tenant isolation or enterprise RBAC.
- Impossible to evaluate deterministically.
- Dependent on unrestricted tool execution.
- Weakly differentiated from existing Stoquify capabilities.
- Legally or operationally unsuitable for Stoquify’s target markets.
- Unclear about data ownership, licensing, or privacy.
- Expensive without a credible path to measurable value.

Phase 3 — Generate and Screen Candidates

Build an initial longlist of at least 30 candidate agents or skills.

For each candidate, determine:

- The user problem it addresses.
- Its primary users and jobs to be done.
- Whether Stoquify already has an equivalent capability.
- Whether it should be built, adapted, partnered for, or rejected.
- Its required data, services, tools, integrations, and permissions.
- Its likely usage frequency.
- Its contribution to retention, referrals, monetization, and defensibility.
- Its security, compliance, and operational risk.

• Its implementation complexity and dependency on the Copilot runtime backbone.

Consolidate overlapping candidates before selecting the final 12.

Do not artificially fill categories. Select the strongest candidates based on evidence and Stoquify-specific value.

Phase 4 — Evaluation Framework

Score every shortlisted candidate from 1 to 10 against the following weighted criteria:

Criterion	Weight
Alignment with Stoquify's core mission	15%
Daily or weekly usage potential	12%
Retention and platform stickiness	12%
Measurable customer value	10%
Referral and collaboration potential	8%
Data, workflow, evidence, or network moat	10%
African and OHADA market relevance	8%
Technical feasibility	8%
Compatibility with the Copilot backbone	5%
Security, privacy, and governance readiness	5%
Monetization or expansion potential	4%
Time to first production value	3%

Calculate the weighted score transparently.

A high score must be supported by concrete evidence and reasoning. Do not manipulate scores to justify a predetermined list.

Required Analysis for Each of the Final 12

For every selected agent or skill, provide:

1. Professional name

• Use a clear, enterprise-appropriate name.

2. Classification

• Agent, skill, shared service, orchestration capability, or hybrid.

3. Problem solved

• Identify the current user pain and why it matters.

4. Target users

• Examples: business owner, CFO, accountant, cashier, inventory manager, payroll administrator, auditor, supplier, employee, or Stoquify operator.

5. Core workflows

• Describe the specific recurring workflows it would execute or support.

6. Features and functionality

• List concrete functions, inputs, outputs, triggers, decisions, and user controls.

7. User experience

• Explain where and when the capability appears in the product and why users would return to it.

8. Daily-use mechanism

• Explain the recurring operational trigger that creates habitual use.

9. Stickiness contribution

• Explain the accumulated value that makes Stoquify increasingly difficult to replace.

10. Referral or collaboration loop

• Explain how it encourages users to invite accountants, employees, suppliers, customers, auditors, or other businesses.

11. Strategic moat

• Identify whether it builds a data, workflow, evidence, integration, trust, compliance, learning, or network moat.

12. Commercial value

• Explain its effect on acquisition, conversion, retention, expansion, pricing power, or cost-to-serve.

13. Pros

• Present specific, Stoquify-relevant advantages.

14. Cons and trade-offs

• Present real limitations, failure modes, dependencies, and opportunity costs.

15. Security and control requirements

• Cover tenant isolation, RBAC, entitlements, redaction, auditability, approvals, human oversight, tool allowlists, evidence requirements, and kill-switch behavior.

16. Data and integration requirements

• Identify required internal services, external systems, APIs, events, and data quality assumptions.

17. Runtime-backbone dependencies

• Identify which Copilot runtime components must exist before it can operate safely.

18. Build, adapt, buy, or partner recommendation

• Give a clear decision with justification.

19. Implementation feasibility

• Rate complexity, estimated effort, critical dependencies, and major risks.

20. Recommended delivery phase

• Now, next, later, experimental, or reject.

21. Success metrics

• Define measurable adoption, retention, efficiency, accuracy, referral, revenue, and risk indicators.

22. Weighted score

• Show the score and explain its strongest and weakest dimensions.

23. Research references

• Link to the most relevant primary sources, implementations, or comparable products.

Portfolio-Level Requirements

After evaluating the final 12, analyze them as a coordinated portfolio rather than 12 isolated tools.

Identify:

• The best three capabilities to implement first.

• The recommended implementation sequence.

• Shared services and infrastructure needed by multiple agents.

• Dependencies and sequencing conflicts.

• Capabilities that should be combined.

• Capabilities that should remain independent for control reasons.

- Where human approval is mandatory.
- Which capabilities can initially operate in read-only or advisory mode.
- Which capabilities could eventually support controlled actions.
- Which capability is most likely to drive daily use.
- Which capability creates the strongest referral loop.
- Which capability creates the strongest defensible moat.
- Which capability offers the fastest commercial return.
- Which capability carries the greatest security or compliance risk.
- Which ideas should explicitly be rejected or postponed.

Required Deliverables

Produce the following sections:

1. Executive Recommendation

Summarize the most important findings, the recommended portfolio, and the top three priorities.

2. Stoquify Current-State Assessment

Document existing agents, skills, architecture, product gaps, user workflows, and relevant constraints.

3. Research Methodology and Sources

Explain where and how the research was conducted, the source-selection standard, and any evidence limitations.

4. Candidate Longlist

Present at least 30 candidates with concise disposition decisions:

- Shortlist
- Merge with another capability
- Later
- Reject
- Already covered

5. Ranked Final 12

Present a comparison table containing:

- Rank
- Agent or skill
- Classification
- Primary user
- Core problem
- Usage frequency
- Stickiness potential

- Referral potential
- Moat potential
- Feasibility
- Risk
- Weighted score
- Delivery phase

6. Detailed Candidate Assessments

Provide the complete required analysis for each of the final 12.

7. Portfolio Architecture

Show how the recommended agents and skills would fit into the Stoquify Copilot backbone, including:

- Copilot entry points
- Agent dispatcher
- Skill and tool registry
- Trusted tenant and authorization context
- Workflow and case orchestration
- Evidence and approval services
- Domain services
- Event infrastructure
- Model gateway
- Audit, telemetry, cost controls, feature flags, suspension, and rollback

Include a Mermaid architecture diagram.

8. Implementation Roadmap

Organize delivery into:

- Foundation prerequisites
- Read-only internal pilot
- Controlled customer pilot
- Expansion
- Advanced and action-capable agents

For each phase, state the prerequisites, deliverables, validation gates, expected value, and explicit non-goals.

9. Measurement Framework

Define leading and lagging metrics for:

- DAU/WAU
- Workflow completion
- Repeat usage
- Time-to-value

- Retention
- Invitations and referrals
- Cross-module adoption
- Expansion revenue
- Decision quality
- Automation accuracy
- User trust
- Approval and rejection rates
- Cost per successful outcome
- Incident and rollback rates

10. Final Decisions

End with direct answers to:

1. Which 12 agents or skills should Stoquify pursue?
2. Which three should be implemented first?
3. What should not be built yet?
4. What infrastructure must exist first?
5. What measurable value should each priority deliver?
6. What are the principal risks?
7. What leadership decisions are required now?

Risk Controls

The recommendations must preserve:

- Tenant isolation.
- Server-side authorization.
- RBAC and module entitlement.
- Least-privilege tool access.
- Exact service and action allowlists.
- Maker-checker separation.
- Fresh authentication for sensitive operations.
- Durable approvals and evidence.
- Full auditability and provenance.
- Privacy and data minimization.
- Safe error handling and redaction.
- Human review for financial, statutory, payroll, compliance, and irreversible actions.
- Feature flags, cost ceilings, kill switches, suspension, rollback, shadow testing, and canary releases.

Never recommend allowing a language model to become the source of truth for accounting, inventory, payroll, statutory, or financial records.

Non-Goals

Do not:

- Implement or install the proposed capabilities during this research task.
- Recommend generic AI chat solely for novelty.
- duplicate existing Stoquify agents or skills.
- Treat vendor claims as independently proven.
- Recommend uncontrolled autonomous financial actions.
- Suggest a premature microservice or agent-swarm architecture.
- Hide important disadvantages or implementation risks.
- assume that internet popularity equals Stoquify relevance.
- Produce a list without repository evidence and comparative scoring.

Success Criteria

The analysis is complete only when:

- Stoquify’s existing capabilities have been inspected.
- At least 30 candidates have been evaluated.
- The final 12 are ranked using the weighted framework.
- Every recommendation includes concrete pros, cons, value, risks, dependencies, and metrics.
- Every material external claim has a source.
- Recommendations are specific to Stoquify and its target markets.
- The top three candidates form a feasible first implementation portfolio.
- Architectural and runtime prerequisites are explicit.
- The report clearly distinguishes verified evidence from inference.
- Leadership can make build, defer, partner, or reject decisions directly from the report.