

Stoquify Top 12 Agents and Skills — Complete Execution Roadmap

Programme date: 2 August 2026
Decision horizon: 30–36 months, with gated value releases from the first 90 days
Authority: Execution roadmap derived from `STOQUIFY_TOP_12_ADDITIVE_AGENTS_AND_SKILLS_RESEARCH_REPORT_2026-08-02.md` and reconciled against the current Stoquify repository
Status: Implementation-ready programme plan; production activation remains fail-closed

1. Executive decision

Stoquify should approve the complete portfolio as a sequenced programme, but should not launch twelve independent agent projects. It should build one governed execution spine, seven shared foundations, and a series of narrow workflow slices that reuse those foundations.

The recommended order is:

- 1. Qualify the existing runtime and close release-evidence blockers.** Product design can continue, but no new external or financially consequential execution should bypass the blocked Phase 2B entry gate.
- 2. Build the shared data-trust, document, consent, case, communication, external-identity, evaluation, and product-measurement foundations.**
- 3. Release the first value loop:** Mobile Document Intake → Receivables and Collections → Advisor Collaboration, supported by Connector Health.
- 4. Add forward decisions and controls:** Cash Forecast, Profitability/Pricing, Continuous Controls, and the shared Omnichannel Communications skill.
- 5. Add network and partner workflows:** Supplier Collaboration and Financing Readiness.
- 6. Expand into revenue and vertical economics:** Sales Pipeline-to-Cash and Job/Project Margin.

The programme is a **conditional GO** for controlled development and a **NO-GO** for broad autonomy or production activation today. The existing runtime has material foundations, but current repository evidence still records blocked production prerequisites, blocked operational evidence, blocked Phase 2B entry, and 999 structurally defined evaluation cases that have not been executed as a certification corpus.

Leadership outcome

Fund the programme in stage gates. Authorize Phase 0 and Phase 1 immediately. Release later funding only when measurable trust, adoption, unit economics, and safety thresholds are met. Keep all messages, accounting mutations, payment actions, credit disclosures, price changes, legal threats, bank-detail changes, and statutory representations human-approved until a separately authorized control proves bounded execution safe.

2. Scope, assumptions, and definitions

2.1 Capabilities in scope

ID	Capability	Portfolio role	Initial autonomy
C01	Receivables and Collections Agent	Daily cash outcome	L2 draft-and-approve
C02	Mobile Document Intake and Evidence Capture	Evidence acquisition foundation	L1 read/extract; L2 verified draft
C03	Advisor Collaboration and Review Agent	Referral and trust network	L1 read-only review; L2 draft requests
C04	Cash Forecast and Scenario Planning Agent	Forward decision support	L1 deterministic insight
C05	Omnichannel Business Communications Skill	Shared action channel	L2 approve-before-send
C06	Continuous Controls and Fraud Signals Agent	Cross-domain loss prevention	L1 confidential signals
C07	Profitability and Pricing Intelligence Agent	Margin decision support	L1 analysis; L2 price draft
C08	Connector Health and Data Trust Agent	Reliability prerequisite	L1 monitoring; bounded operator remediation only
C09	Supplier Collaboration Network Agent	Counterparty network	L1/L2 scoped collaboration
C10	Financing Readiness and Consent Package Agent	Evidence-to-finance partner route	L2 consented package preparation
C11	Sales Pipeline-to-Cash Agent	Revenue workflow	L2 quote/follow-up drafts
C12	Job and Project Margin Agent	Vertical economics	L1 analysis; L2 estimate/change drafts

Autonomy levels. L0 is disabled. L1 is read-only analysis. L2 prepares a reviewable draft and requires an authorized human action. L3 is a bounded, reversible, policy-controlled action with an explicit service allowlist and approval policy. L4 autonomous consequential execution is outside this roadmap.

2.2 Planning assumptions

- The existing tenant, RBAC, ledger, evidence, provider-event, agent-run, release-control, and Command Agent foundations will be extended rather than replaced.
- Service-owned deterministic code remains the source of truth for amounts, eligibility, accounting, permissions, workflow state, and compliance decisions. Models may classify, summarize, draft, and explain; they do not invent business truth.
- Country coverage is enabled through versioned country packs and legal review, not through a universal “Africa” rule set.
- Estimates are planning ranges, not vendor quotations. Cost ranges assume loaded delivery cost of **€7,000–€13,000 per FTE-month**, plus **€5,000–€35,000 per month** in pilot-stage model, messaging, document-processing, observability, and connector costs.
- External prerequisites—production targets, secrets, workload identity, approvers, statutory experts, provider credentials, pilot tenants, and operational owners—are calendar risks outside engineering’s unilateral control.

3. Evidence-led current-state reconciliation

3.1 What is present

Area	Status	Repository evidence	Roadmap consequence
Agent definitions	PRESENT	9 agents, 28 skills, 37 registry capabilities validate structurally	Extend the registry; do not create disconnected prompts
Evaluation definitions	PARTIALLY PRESENT	999 cases validate structurally but remain <code>NOT_TESTED</code>	Convert into executable, evidence-producing certification suites
Tenant-scoped runtime	PRESENT, NOT PRODUCTION-QUALIFIED	<code>AgentRun</code> , <code>steps</code> , <code>evidence</code> , <code>feedback</code> , <code>cost</code> and <code>incidents</code> ; <code>runner/context/policy</code> services	Reuse, then qualify through existing gates
Tool registry and policy	PRESENT	Exact tool definitions, policy checks, prohibited-pattern gates	Add capability-specific tools only through allowlisted contracts
Rollout controls	PRESENT	Shadow mode, UI visibility control, suspension and kill switch	Preserve fail-closed rollout pattern
Command Agent slice	PRESENT, INACTIVE/CONTROLLED	Governed <code>service/action/UI</code> path and daily digest integration	Use as base-runtime pilot before Top 12 execution
Provider-event evidence	PRESENT	Provider <code>accounts/events</code> , <code>signatures</code> , <code>idempotency</code> and <code>correlation</code>	Extend to common connector-health contracts
Generic external collaboration	PARTIALLY PRESENT	Tenant invitation exists; purpose-bound firm/client/counterparty access is not established	Build external-party identity and resource scopes
Durable generic workflow cases	PARTIALLY PRESENT	Agent runs/steps/proposals exist; capability-neutral case lifecycle is not clearly established	Add a case aggregate or extend an existing service-owned case contract
General consent ledger	PARTIALLY PRESENT	Domain consent exists; cross-channel, purpose-bound, revocable consent is not proven	Build a common consent service before messaging or finance sharing
Document intake	PARTIALLY PRESENT	Upload/evidence patterns exist; universal malware-scanned ingestion, extraction provenance and verification are not proven	Build a reusable intake pipeline
Model-provider runtime	NOT PROVEN ACTIVE	Current release evidence identifies <code>modelProvider: none</code> ; no provider SDK is established in the reviewed package manifest	Introduce provider abstraction only after privacy, cost and evaluation contracts

3.2 Qualification evidence and limitations


- The July runtime completion audit records Phase 0/1 and a narrow inactive Phase 2A slice as implemented, but marks promotion before Phase 2B as blocked.
- The July blocker-elimination report records `REJECTED_NO_GO` for pilot/production while allowing fail-closed development to continue.
- Operational release evidence records **152 blockers**; external-input readiness records **1 of 13** prerequisite groups ready and **102 blockers**.

- Production database targeting, secrets, workload identity, statutory authority, provider credentials, operational ownership, freeze reconciliation, governance approval, Phase 2B entry, and Phase 3 authority remain material gates.
- The architecture graph is dated 14 June 2026 and contains 4,121 nodes, 5,321 edges and 135 communities. It confirms tenant defence, ledger-first operational posting, RBAC hardening, and error-handling control patterns, but it predates the July runtime work and is used as navigation—not current release proof.

3.3 Non-negotiable implication

The Top 12 programme may perform discovery, contract design, fixtures, deterministic calculation work, UX prototypes and internal shadow evaluation while external release blockers are being closed. It may not treat those activities as production authorization. Every stage must emit independent evidence and must inherit the stricter of the base-runtime gate and the capability-specific gate.

4. Target operating architecture



Architectural invariants

1. Organization and actor context comes from authenticated server state, never model or client claims.
2. Every tool has a stable ID, versioned input/output schema, authorization policy, idempotency behavior and service owner.
3. Financial and compliance calculations are deterministic and independently testable.
4. Every material output links to source evidence, freshness, assumptions and transformation version.
5. Every side effect has preview, explicit authority, immutable audit, replay protection and safe retry semantics.
6. External access is purpose-bound, resource-bound, time-bound, revocable and denied by default.
7. Model prompts and outputs never receive raw secrets and receive only the minimum necessary personal data.
8. Feature flags, tenant allowlists and kill switches apply at agent, skill, tool, provider, country and tenant levels.
9. A capability cannot promote if the base runtime, its dependencies, or its own evaluation gate is red.

5. Shared foundation workstreams

ID	Foundation	Key deliverables	Primary owner	Exit gate	Estimate
F- Runtime and A orchestration	Security and B authorization	Versioned registry loader, dispatcher, durable run/case state, checkpoints, retries, idempotency, resumable approvals	Principal platform engineer	Crash/retry/replay tests; no duplicate side effects	24–40 FTE-weeks
F- Security and B authorization		Trusted tenant/actor context, entitlement service, resource scopes, external identities, maker-checker, fresh auth, purpose-bound consent	Security architect	Cross-tenant and privilege-escalation suite 100% pass	28–48 FTE-weeks
F- Evidence and data C trust	Model controls D	Evidence envelope, hashes, source coordinates, freshness, provenance, document pipeline, connector health and data-quality contracts	Data/platform lead	Every surfaced claim has evidence state; stale data blocks unsafe output	32–54 FTE-weeks
F- Model controls D		Provider abstraction, redaction, prompt/version registry, structured output, deterministic fallback, cost/latency budgets	AI platform lead	Provider failure and prompt-injection tests; budget enforcement	20–36 FTE-weeks
F- Evaluation and E assurance	User experience F	Executable catalog, fixtures, adversarial suites, regression baselines, shadow/canary evidence, independent gate	QA/assurance lead	Required cases executed with signed artifacts and zero critical violations	36–60 FTE-weeks
F- User experience F		Command surfaces, verification inbox, evidence drawer, approval diff, case timeline, offline/mobile states, accessibility	Product/design lead	Task completion and accessibility thresholds met	28–46 FTE-weeks
F- Growth and G commercialization		Packaging, entitlements, usage meters, invitation attribution, value telemetry, experiments, customer-success playbooks	Product growth lead	Measured activation/value loop and auditable billing events	18–30 FTE-weeks

Foundation implementation notes


- F-A — runtime.** Extend the current `AgentRun` and `step/proposal` structures. First publish a decision record on whether a generic workflow-case aggregate is required. Do not force long-lived collection, advisor, supplier or investigation cases into short-lived run semantics if their lifecycle, ownership and legal retention differ.
- F-B — security.** Add `ExternalParty`, `OrganizationRelationship`, `ResourceGrant`, `PurposeConsent`, `ConsentEvent`, `AccessSession` and revocation propagation concepts, or equivalent service-owned schemas. Establish row-level authorization tests for client-firm, buyer-supplier, business-customer and business-finance-partner relationships.
- F-C — evidence.** Standardize an evidence envelope containing organization, source, source version, captured-at, observed-at, freshness policy, hash, transformation chain, confidence, reviewer, retention class and redaction state. Build document ingestion as an asynchronous pipeline: quarantine → scan → identify → extract → validate → duplicate check → human verify → service-owned draft.
- F-D — models.** Start with a provider-neutral adapter. Route low-risk classification and drafting to the least expensive qualified model. Use deterministic or template fallbacks for critical queues. Store prompt/version hashes and evaluation cohort—not unrestricted chain-of-thought.
- F-E — assurance.** Convert 999 cases from catalog entries into runnable fixtures with expected policy, tool, evidence and output assertions. Add cases for every new capability; maintain critical suites for cross-tenant denial, prompt injection, stale evidence, side-effect approval, replay, provider outage, PII leakage and country-pack expiry.
- F-F — UX.** The primary interface is a queue or business cockpit, not a blank chat box. Every recommendation must show “why,” sources, freshness, uncertainty, what will happen, who must approve and how to undo or contest it.
- F-G — growth.** Instrument time-to-first-value, weekly completed outcomes, invitations, accepted invitations, multi-user activation, retained value, expansion, cost-to-serve and trust incidents. Never optimize message volume or invitations without recipient consent and complaint safeguards.

6. Dependency map and critical path

6.1 Capability dependencies

Capability	Hard prerequisites	Soft prerequisites	Unlocks
C08 Connector Health	F-A, F-B, F-C, provider contracts	F-E	Every data-dependent capability
C02 Document Intake	F-A–F-F, object storage/security	C08	C01, C03, C06, C09, C10, C12
C01 Collections	AR truth, C08, cases, consent, approval	C02, C05	C04, C10, C11
C03 Advisor Collaboration	External identity, resource grants, audit	C02, C05	C09, C10, advisor referral loop
C05 Communications	Consent, cases, provider allowlists, webhook verification	C08	C01, C03, C09, C11
C04 Cash Forecast	Trusted cash/AR/AP/payroll read models, C08	C01, C02	C10, executive planning
C07 Profitability/Pricing	Cost/revenue read models, deterministic allocation engine	C02, C08	C11, C12
C06 Controls/Fraud	Normalized events, confidential cases, F-E	C02, C08	Assurance tier
C09 Supplier Network	Proven C03 external access, C02, C05	C06	Counterparty network
C10 Financing Readiness	Mature evidence grading, consent, C03, C04	C01, C02	Regulated partner channel
C11 Pipeline-to-Cash	C01, C05, C07, quote/order services	C03	Revenue suite
C12 Job/Project Margin	C07, projects/time/cost models	C02, C03, C11	Vertical packages

6.2 Critical path



Calendar-critical external inputs are the production deployment target, CI candidate identity, secrets/workload identity, statutory reviewers, provider credentials, named product/security/operations owners, pilot cohort and activation authority. While those are pending, teams may complete contracts, fixtures, synthetic tests and design prototypes; they may not self-authorize a pilot.

7. Phased delivery roadmap

Phase	Calendar	Outcomes	Promotion gate	Indicative cost
0. Mobilize and reconcile	Weeks 0–8	Programme charter; owner map; authoritative status register; current gate rerun; baseline metrics; architecture decisions; top-three discovery	Evidence register approved; no status contradictions	€250k–€600k
1. Qualify the spine	Months 2–6	Close base-runtime blockers; executable evaluation harness; consent/external identity/case/evidence contracts; connector health MVP	Phase 2B base-runtime entry 100%; critical security tests pass	€700k–€1.3m
2. First value loop	Months 5–11	Document Intake pilot; draft-only Collections; read-only Advisor workspace; email communication pilot	Each pilot meets safety, accuracy, task-value and rollback gates	€900k–€1.6m
3. Decisions and controls	Months 9–16	13-week cash forecast; profitability; narrow fraud-signal pilots; one approved mobile messaging channel	Deterministic accuracy, false-positive and consent thresholds	€800k–€1.5m
4. Network workflows	Months 14–22	Advisor firm cockpit; Supplier Collaboration; consented Financing Readiness with one regulated partner	External access, legal, fairness and partner assurance	€900k–€1.7m
5. Revenue and verticals	Months 20–30	Narrow Pipeline-to-Cash; one Job/Project Margin vertical	Quote/margin/project data-quality and commercial proof	€900k–€1.8m
6. Scale and optimize	Month 24 onward	Country/connector expansion, L3 candidates, marketplace/partner operations, cost optimization	Per-country and per-tool independent promotion	€400k–€1.0m/year incremental

Phases overlap only where dependencies permit. A recommended 16–22-person cross-functional programme can target 24–30 months. A minimum 10–13-person team should plan for 30–42 months and fewer concurrent pilots. Total internal planning envelope is **€4.8m–€7.8m over 24–30 months**, excluding broad go-to-market spend, acquisitions, regulated-partner capital and country-specific licensing. Funding should be released in gate-based tranches rather than committed as one irreversible amount.

8. Capability execution playbooks

C01 — Receivables and Collections Agent

- **Outcome and users:** Give owners, CFOs and AR staff a daily “get paid” queue that converts overdue invoices into evidence-linked, approved action.
- **MVP:** Aging/risk queue; dispute suppression; contact and consent validation; explainable priority; editable email draft; approval; delivery webhook; promise-to-pay case; reconciled-payment closure.
- **Inputs/outputs:** Customer/invoice/payment history, dispute flags and channel permissions → ranked case, reason codes, message draft, promise and next action.
- **Build/buy:** Build workflow, policy and outcome graph; buy regulated message transport.
- **Controls:** No auto-send; no legal threat, discount, fee waiver, account hold or restructuring without explicit authority; quiet hours, suppression and complaint controls.
- **Milestones:** data contract (2 weeks); queue and rules (4); draft/approval (4); delivery/replies (3); shadow and pilot (4–6). **Estimate:** 12–19 weeks, 5–7 FTE.
- **Acceptance:** zero cross-tenant contacts; zero duplicate sends; 100% evidence-linked priority; ≥90% delivery-state reconciliation; complaint/opt-out below approved threshold.
- **Value hypothesis:** reduce median days-to-pay by 5–12% in eligible pilot invoices and create completed collection work on ≥3 days/week for activated AR users.

C02 — Mobile Document Intake and Evidence Capture

- **Outcome and users:** Turn photos, PDFs and attachments into verified service-owned drafts without losing provenance.
- **MVP:** Mobile/web upload; signed URL; quarantine and malware scan; receipt/supplier-invoice classification; field extraction; source highlighting; confidence policy; duplicate detection; verification inbox; AP/evidence draft handoff.
- **Controls:** No silent posting; encrypted storage; tenant keys; retention class; PII redaction; content/size checks; sensitive-document reviewer separation.
- **Build/buy:** Buy/adapt OCR and scanning components; build provenance, validation, country layouts and service boundaries.
- **Milestones:** threat model/storage (3); pipeline (5); extraction adapters (5); reviewer UX (4); offline/retry (3); pilot (4–6). **Estimate:** 16–24 weeks, 6–9 FTE.
- **Acceptance:** ≥95% verified accuracy for required header fields in supported document classes; 100% originals/hash/provenance retained; no malware escapes; documented unit cost.
- **Value hypothesis:** cut receipt-to-verified-draft median time by ≥50%; reduce missing-evidence exceptions by ≥20% in activated tenants.

C03 — Advisor Collaboration and Review Agent

- **Outcome and users:** Give accountants and advisors a purpose-limited multi-client review cockpit and give businesses a simple request queue.
- **MVP:** advisor invitation; firm membership; client-granted scope; evidence request; threaded response; read-only exception/report access; expiry/revocation; access audit.
- **Controls:** no cross-client discovery; resource and purpose scopes; fresh auth; watermark/download controls; explicit advisory/reviewed/certified states.
- **Build/buy:** Build Stoguify trust graph and workspace; buy only commodity identity verification if required.
- **Milestones:** authorization model (4); invitation/relationship (4); request case (4); cockpit (5); denial/adversarial tests (3); pilot (4–6). **Estimate:** 18–26 weeks, 6–8 FTE.
- **Acceptance:** 100% cross-client denial suite; revocation effective within defined SLA; full access audit; ≥60% pilot invitation acceptance.
- **Value hypothesis:** one activated advisor brings or manages ≥3 client organizations over time; reduce evidence-request cycle time by ≥30%.

C04 — Cash Forecast and Scenario Planning Agent

- **Outcome:** Deterministic 13-week base/downside/upside cash scenarios with explicit assumptions and confidence—not fabricated precision.
- **MVP:** opening cash, AR probability, AP commitments, payroll/tax and recurring flows; assumption ledger; scenario comparison; freshness and driver waterfall; Command card.
- **Controls:** booked facts and assumptions separated visually and in schema; no automated transaction; forecast blocked or degraded when evidence is stale.
- **Build/buy:** Build calculation and scenario storage; optionally use a statistical library for calibrated baselines.
- **Estimate/gate:** 14–22 weeks, 5–7 FTE; pass backtests by horizon and achieve 100% source/assumption traceability.
- **Value hypothesis:** ≥40% of activated owner/CFO users review weekly; forecast error improves against naive baseline over two cycles.

C05 — Omnichannel Business Communications Skill

- **Outcome:** A reusable, consent-aware message pipeline for all agents.
- **MVP:** versioned templates; preview/edit/approve; email plus one country/provider pilot channel; verified webhooks; inbound reply linkage; opt-out; rate limiting; delivery reconciliation.
- **Controls:** exact provider allowlists; secrets isolation; replay protection; sender verification; quiet hours; emergency suspension; no model access to credentials.
- **Build/buy:** Build policy/orchestration; buy transports.
- **Estimate/gate:** 12–20 weeks, 5–7 FTE; zero duplicate sends in fault injection, 100% consent checks, provider failure recovery demonstrated.
- **Value hypothesis:** reduce workflow response time ≥25% without breaching complaint and opt-out thresholds.

C06 — Continuous Controls and Fraud Signals Agent

- **Outcome:** Surface confidential, reviewable signals across AP, cash, stock, payroll, approvals and access—not accusations.
- **MVP:** supplier bank-detail change, duplicate-payment/document, approval bypass and payment-destination collision rules; materiality; evidence pack; restricted case; outcome label.
- **Controls:** deterministic features; “signal/requires review” wording; restricted access; no automated discipline or payment block protected investigation records.
- **Build/buy:** Build control graph and cases; use statistical libraries only after deterministic baseline.
- **Estimate/gate:** 16–26 weeks, 6–9 FTE; confirmed-signal precision target set per control, false-positive budget, zero confidential-case leakage.
- **Value hypothesis:** reduce control review time and establish measurable prevented-loss evidence without alert fatigue.

C07 — Profitability and Pricing Intelligence Agent

- **Outcome:** Explain margin by item, customer, branch and channel and prepare controlled scenarios.
- **MVP:** contribution-margin engine; cost freshness; versioned allocations; discount leakage; mix/cost waterfall; scenario sandbox; price-floor draft.
- **Controls:** deterministic calculations; assumptions visible; role-scoped margins; no automatic published price changes.
- **Build/buy:** Build native unit-economics layer.
- **Estimate/gate:** 14–22 weeks, 5–7 FTE; calculations reconcile to service-owned totals within agreed tolerance and all allocations are reproducible.
- **Value hypothesis:** identify actionable leakage in ≥30% of eligible pilots and measure realized improvement after approved actions.

C08 — Connector Health and Data Trust Agent

- **Outcome:** Make stale, missing, duplicated or broken data visible before it contaminates an agent decision.
- **MVP:** connector inventory; last success; SLA/freshness; gap/duplicate/schema drift; credential-expiry metadata; affected-capability map; dead-letter and incident linkage; safe reconnect guidance.
- **Controls:** credentials never enter prompts/logs; signed webhooks; operator approval for credential change; provider isolation.
- **Build/buy:** Build common trust plane; partner country by country for data access.
- **Estimate/gate:** 12–18 weeks, 5–7 FTE; detection/recovery objectives met and stale-dependent agent outputs fail closed.
- **Value hypothesis:** reduce connector-caused support tickets ≥20% and mean time to detection ≥50% in instrumented connectors.

C09 — Supplier Collaboration Network Agent

- **Outcome:** Move PO acknowledgement, evidence, delivery commitments and disputes into scoped supplier cases.
- **MVP:** supplier invite/signed link; PO acknowledgement; delivery commitment; attachment upload; invoice dispute thread; limited payment status.
- **Controls:** PO/invoice-specific grants; no tenant browsing; attachment scan; fresh auth and independent maker-checker for bank-detail changes.
- **Build/buy:** Build after C03 proves external authorization.
- **Estimate/gate:** 18–28 weeks, 6–9 FTE; external denial suite, bank-change controls, dispute audit and supplier adoption thresholds pass.
- **Value hypothesis:** reduce acknowledgement and dispute cycle times ≥25%; create consented supplier-to-tenant invitations.

C10 — Financing Readiness and Consent Package Agent

- **Outcome:** Convert daily trusted operations into a consented, time-limited finance evidence package; never claim credit approval.
- **MVP:** readiness checklist; evidence gaps; reconciled summaries; redaction; granular consent; package hash/version; expiry/revocation; partner access log; correction process.
- **Controls:** regulated-partner diligence; country legal review; purpose limitation; no protected-attribute inference, hidden score or data sale; fairness and complaint review.
- **Build/buy:** Build readiness/package layer; partner with regulated institutions.
- **Estimate/gate:** 22–36 weeks plus partner/legal lead time, 6–10 FTE; legal, privacy, fairness, revocation and partner-security gates all green.
- **Value hypothesis:** reduce evidence-preparation time ≥40% and partner turnaround without increasing complaints or unauthorized access.

C11 — Sales Pipeline-to-Cash Agent

- **Outcome:** Connect a narrow opportunity/quote workflow to approved prices, order/invoice creation, collections and margin.
- **MVP:** lead inbox; qualification; quote draft; margin/discount policy; approved follow-up; acceptance evidence; service-owned invoice handoff; win/loss reason.
- **Controls:** no generic CRM sprawl; contact consent; approved price/discount boundaries; no unapproved commitment or message.
- **Build/buy:** Build narrow native workflow; integrate rather than replace full CRMs when needed.
- **Estimate/gate:** 18–28 weeks, 6–9 FTE; quote correctness, authorization, margin and conversion telemetry pass.
- **Value hypothesis:** reduce quote preparation and quote-to-cash time ≥25% in the target segment.

C12 — Job and Project Margin Agent

- **Outcome:** Protect margin for one chosen vertical through estimate, commitment, time/cost, change and billing readiness.
- **MVP:** cost codes; estimate draft; budget/actual/commitment; time/expense inputs; change evidence; forecast-at-completion; billing checklist.
- **Controls:** project role scopes; versioned assumptions; client/subcontractor grants; no automatic contract, change order or invoice release.
- **Build/buy:** Build as a vertical package after C07 and C11; select one segment through discovery.
- **Estimate/gate:** 20–32 weeks, 7–10 FTE; source completeness, reproducible forecast, access and vertical product-market proof pass.
- **Value hypothesis:** reduce unbilled work and detect material margin drift earlier in ≥50% of at-risk pilot projects.

9. Integrated work breakdown structure

The following epics are issue-tracker ready. “Done evidence” is mandatory, not optional documentation.

WBS	Epic	Depends on	Owner	Done evidence
PGM-001	Charter, scope, autonomy policy and funding gates—		Executive sponsor	Signed charter and decision log
PGM-002	Authoritative programme status register	PGM-001	Programme lead	Reconciled source-linked register
PGM-003	Named RACI, on-call and approvers	PGM-001	COO/CTO	Accepted ownership roster
PGM-004	Benefits baseline and pilot cohorts	PGM-001	Product/data	Baseline dashboard and cohort consent
RUN-001	Rerun base runtime entry gates	PGM-002	Platform assurance	Machine-readable gate artifacts
RUN-002	Close production DB/target and artifact identity	RUN-001	Platform/DevOps	Immutable candidate proof
RUN-003	Secrets, workload identity and rotation	RUN-002	Security/DevOps	Secret scan and rotation drill
RUN-004	Operational evidence collectors and alerts	RUN-002	SRE	Successful dry run and paging evidence
RUN-005	Freeze/drift reconciliation	RUN-001	Release manager	Zero unexplained drift
RUN-006	Governance and independent approval	RUN-003, RUN-004, RUN-005	Assurance	Signed gate decision
RUN-007	Phase 2B base Command pilot	RUN-006	Product/platform	Pilot report and rollback drill
FND-001	Case/run lifecycle ADR	RUN-001	Architect	Approved ADR and schema decision
FND-002	Durable case/checkpoint implementation	FND-001	Platform	Recovery/idempotency tests
FND-003	External-party identity and relationship model	RUN-003	Identity lead	Threat model and denial suite
FND-004	Purpose/resource/time-bound grants	FND-003	Security	Cross-tenant tests 100% pass
FND-005	General consent ledger and revocation	FND-003	Privacy lead	Consent lifecycle tests
FND-006	Evidence envelope and freshness policy	—	Data trust lead	Contract tests and sample lineage
FND-007	Document quarantine/storage/scanning	RUN-003, FND-006	Platform/security	Malware and retention tests
FND-008	Provider-neutral model router	RUN-006	AI platform	Failure, budget and redaction tests
FND-009	Structured-output and prompt registry	FND-008	AI platform	Version and schema tests
FND-010	Executable evaluation harness	RUN-001	QA	Reproducible local/CI run
FND-011	Convert 999 catalog cases	FND-010	QA/domain owners	999 executed case artifacts
FND-012	Common connector health contract	FND-006	Integration lead	Adapter contract suite
FND-013	Evidence/approval/case UX components	FND-002, FND-006	Design/frontend	Accessibility and usability evidence
C08-001	Connector inventory and freshness read model	FND-012	Integration	Reconciliation tests
C08-002	Gap/drift/credential warning rules	C08-001	Integration/SRE	Fault-injection results
C08-003	Trust banners and remediation queue	C08-002, FND-013	Product	Task test and telemetry
C02-001	Receipt/invoice classification and extraction	FND-007, FND-009	Document team	Golden-set accuracy report
C02-002	Duplicate and confidence policy	C02-001	Data trust	Adversarial fixture results
C02-003	Verification inbox and AP draft handoff	C02-002, FND-013	Product/AP	End-to-end evidence chain
C05-001	Template, consent and provider contracts	FND-005, FND-012	Communications	Policy tests
C05-002	Approve/send/reconcile email flow	C05-001, FND-002	Communications	Retry/replay/delivery tests
C05-003	One mobile-message pilot adapter	C05-002	Communications	Provider/legal/pilot gate
C01-001	AR risk and priority read model	C08-003	AR/data	Deterministic fixture tests
C01-002	Collection case and dispute suppression	C01-001, FND-002	AR	Workflow tests
C01-003	Draft, approval, promise and closure loop	C01-002, C05-002	AR/product	Shadow and pilot outcome report
C03-001	Advisor invitation and firm relationship	FND-004	Identity/product	Access test suite
C03-002	Evidence request and response case	C03-001, C02-003	Advisor team	End-to-end pilot evidence
C03-003	Multi-client advisor cockpit	C03-002	Advisor team	No cross-client leakage; usability pass

C04-001	Forecast source and assumption contracts	C08-003	Finance/data	Reconciled sample dataset
C04-002	Deterministic 13-week engine	C04-001	Finance engineering	Backtest report
C04-003	Scenario UX and Command integration	C04-002	Product	User comprehension test
C07-001	Contribution-margin and allocation engine	C08-003	Finance engineering	Ledger/read-model reconciliation
C07-002	Scenario and price-draft controls	C07-001	Product	Approval/assumption tests
C06-001	Confidential investigation case/access	FND-004, FND-002	Security	Restricted-access tests
C06-002	Bank-change and duplicate-payment signals	C06-001, C02-002	Controls	Precision/false-positive report
C06-003	Cross-domain signal outcome loop	C06-002	Controls/data	Outcome and drift dashboard
C09-001	Supplier-specific grants and portal	C03-003	Supplier team	External denial suite
C09-002	PO acknowledgement and evidence	C09-001, C02-003	Supplier/AP	Pilot cycle-time report
C09-003	Disputes and bank-change maker-checker	C09-002, C06-002	Supplier/security	Fraud-control simulation
C10-001	Finance-readiness schema and checklist	C04-003, C03-003	Finance partnerships	Domain review
C10-002	Consent, package, redaction and revocation	C10-001, FND-005	Privacy/security	Privacy and revocation test
C10-003	Regulated partner pilot	C10-002	Partnerships/legal	Signed assurance and outcome report
C11-001	Narrow opportunity and quote contracts	C07-002	Revenue product	Scope and domain acceptance
C11-002	Quote approval and communication loop	C11-001, C05-003	Revenue team	Price/discount authorization tests
C11-003	Acceptance-to-invoice/collection handoff	C11-002, C01-003	Revenue/AR	End-to-end cash trace
C12-001	Select vertical and define project read model	C07-002, C11-001	Vertical GM	Signed discovery evidence
C12-002	Estimate, cost and change controls	C12-001, C02-003	Vertical team	Golden project scenarios
C12-003	Forecast-at-completion and billing readiness	C12-002, C11-003	Vertical team	Pilot margin report
REL-001	Per-capability shadow and canary framework	FND-010	Assurance	Promotion/rollback proof
REL-002	Country-pack and provider certification	REL-001	Legal/integration	Versioned approvals
REL-003	Cost, latency, quality and incident SLOs	FND-008	SRE/FinOps	Operational dashboard
REL-004	Quarterly portfolio continue/pivot/stop review	PGM-004	Steering committee	Decision and benefit evidence

10. Evaluation and release strategy

10.1 Evaluation pyramid

1. **Contract tests:** schemas, tenant context, RBAC, entitlements, consent, tool allowlists and service ownership.
2. **Deterministic unit/golden tests:** calculations, rules, freshness, duplicate detection, idempotency and redaction.
3. **Agent scenario tests:** prompt injection, missing evidence, ambiguity, unsafe requests, refusal and bounded clarification.
4. **Workflow tests:** crash/retry, approval pause/resume, webhook replay, duplicate delivery, revocation, provider outage and rollback.
5. **Security/privacy tests:** cross-tenant and cross-client access, PII leakage, secret exposure, signed-link misuse and malicious documents.
6. **Shadow evaluation:** compare outputs with real workflows without user-visible execution.
7. **Canary pilot:** tenant allowlist, bounded workload, staffed monitoring, predeclared stop conditions.
8. **Independent certification:** Platform Assurance cannot self-certify the capability owner's assertions; an independent release authority signs the promotion artifact.

10.2 999-case conversion plan

- Freeze a catalog version and map every case to owner, risk tier, fixture, expected policy result, expected tool set, evidence requirements and output assertions.
- Implement the highest-risk 150 cases first: tenancy, authorization, prohibited action, stale evidence, financial mutation, external message, consent, privacy, replay and kill switch.
- Execute all 999 in CI against a deterministic stubbed provider and then a qualified live-provider shadow environment where applicable.
- Store run ID, code/artifact identity, model/prompt/tool versions, dataset hash, result, latency, cost and reviewer decision.
- Add capability-specific cases before each pilot; a regression cannot be waived silently.

10.3 Promotion gates

Gate	Minimum evidence
Design ready	Product brief, threat model, data map, tool contracts, metrics, stop conditions
Build complete	Code review, schema/migration proof, deterministic tests, runbook, telemetry
Shadow ready	Critical evaluation pass, no unresolved severity-1/2 issue, cost/latency budget
Canary ready	Independent approval, tenant consent, named on-call, rollback rehearsal, provider/legal clearance
Scale ready	Value threshold met, trust incidents within budget, unit economics viable, support ready
L3 candidate	Separately authorized action, reversible/idempotent tool, maker-checker policy, extended clean canary

11. Value-realization traceability

	Capability	Daily-use value	Retention value	Referral value	Revenue value	Trust/risk value	Defensible asset
C01	Collections	Daily get-paid queue	Outcome history and promises	Customer links and advisor visibility	Premium cash tier/usage	Controlled, evidenced contact	Payment-behavior outcome graph
C02	Document Intake	Snap/upload and verify	Evidence archive and less re-entry	Advisor evidence requests	Usage tiers	Provenance, accuracy, malware controls	Labeled regional document corpus
C03	Advisor	Request/review queue	Multi-user workflow lock-in	Highest—advisor brings clients	Firm edition/client bundles	Scoped access and review states	Firm-client trust graph
C04	Forecast	Runway/driver card	Assumption and outcome history	Share with advisor/lender	Planning tier	Avoids unsupported precision	Forecast/outcome dataset
C05	Communications	Act inside workflow	Interaction continuity	Recipient-to-portal invitations	Usage/premium automation	Consent and immutable messaging	Response/outcome history
C06	Controls	Confidential signal queue	Loss-prevention dependency	Auditor collaboration	Assurance tier	Material control improvement	Cross-domain control graph
C07	Profitability	Margin-move decisions	Proprietary unit economics	Advisor collaboration	Premium analytics	Reproducible assumptions	Business economics history
C08	Connector Health	Background trust/readiness	Reliability of every feature	Limited direct	Enterprise reliability	Prevents stale/false answers	Provider remediation knowledge
C09	Supplier	PO/dispute queue	Counterparty workflow	Supplier network loop	Network/premium services later	Safer bank changes/disputes	Supplier performance graph
C10	Financing	Readiness progression	Operations-to-evidence dependency	Lender/advisor route	Transparent partner/premium model	Consent, redaction, fairness	Consented evidence packages
C11	Pipeline-to-Cash	Revenue next actions	Lead-to-cash history	Quote/customer links	Revenue suite	Price/commitment control	Conversion and cash graph
C12	Project Margin	Project-risk queue	Vertical system of record	Client/subcontractor links	Vertical package	Change/billing evidence	Estimate-to-realized-margin data

11.1 Explicit value mapping matrix

	Capability	Daily-use v alue	Retention value	Referral value	Revenue value	Trust value	Cost-reduction value	Data value	Workflow value	Network v alue	OHADA/African-market value
C01	Collections	Faster daily cash recovery decisions	Reusable payer/contact history and promise outcomes	Referral through advisor/customer visibility	Premium for collections stack, upsell probability	Evidence-backed outreach reduces disputes and pressure	Reduced manual follow-up cost and recall time	Enrichment of payment pattern and customer behavior datasets	Queue-first collection workflows reduce chum	Opens advisor-led customer concentration effects	Aligns with OHADA proof-of-cash evidence requirements
C02	Document Intake	Immediate document capture and extraction in first workflow step	Lower re-entry and filing friction	Advisor/client access to structured evidence	Usage-based intake and add-on evidence bundle revenue	Document integrity and extraction transparency reduce fraud disputes	Saves clerk labor and accelerates audit-ready closure	Builds trusted regional document corpus with redaction/privacy controls	Standardizes intake workflow across channels	Enables supplier/client evidence exchange later	Supports OHADA invoice/supporting-document compliance

C03 Advisor	Shared queue for review + action	Multi-user for stickiness and account entanglement	Advisor referrals from proof-ready workflows	Advisory-tier packs and retainer retention	Scoped permissions, provenance and override transparency	Avoids duplicate calls and untracked advice	Produces advisor interaction evidence graph	Institutionalizes review workflow and approvals	Extends advisor recommendation network	Fits advisor-led close standards in OHADA contexts
C04 Forecast	Daily scenario visibility before decisions	Better planning continuity in volatile seasons	Shareable scenario packages for teams	Planning-tier upgrade and advisory confidence	Guardrails against unsupported precision	Reduces expensive stock/credit misalignment	Strengthens cash/AR/AP dataset quality	Creates recurring planning cadence	Foundation for commercial negotiation prep with suppliers	Supports liquidity control used in OHADA insolvency controls
C05 Communications	In-workflow message control points	Continual operational thread continuity	Partner portal engagement loops	Communication premium and automations	Consent-first trace prevents misuse	Fewer failed follow-ups and lower support churn	Communication outcome logs improve evidence quality	Improves multi-touch sales/cash workflows	Supplier/customer message infrastructure for network effects	Local communication compliance with localized consent
C06 Controls	Continuous risk signal loop in daily ops	Builds trust in controls and process discipline	Shared risk signal view for advisors and auditors	Assurance-tier pricing and compliance moat	Hardens early warning and manual escalation	Prevents losses and late remediations	Creates labeled exception dataset	Embeds controls into process rather than post-facto cleanup	Facilitates supplier and partner trust if integrated	Supports OHADA control evidence retention requirements
C07 Profitability	Instant margin and pricing context in pricing loops	Keeps teams using one source of truth for price discipline	Advisor and client confidence in profitability discussions	Premium analytics bundle, margin visibility upsell	Deterministic contribution logic reduces pricing disputes	Prevents margin erosion and unnecessary discounts	Creates unit-level profitability knowledge base	Standardizes change-request and pricing workflows	Basis for sector templates and cross-tenant benchmarking	Consistent cost accounting under OHADA costing principles
C08 Connector Health	Daily connector trust status in control cockpit	Reduces uncertainty across all daily operations	Faster trustworthy partner onboarding for connected services	Enterprise/regulatory reliability package value	Prevents stale data and false assertions	Reduces reconciliation and operational incident costs	Produces source-quality provenance and freshness telemetry	Unblocks safe multi-module operations	Shared reliability signal supports broader integrations	Demonstrates infrastructure quality governance for host jurisdictions
C09 Supplier	Faster PO/dispute handling in daily commerce	Strengthens supplier relationship continuity	Supplier-network participation and network response referrals	New supplier-network premium later	Consented exchange + role boundaries reduce leakage	Reduces order errors and delays	Builds PO/fulfillment/settlement quality history	Centralizes supplier workflows and commitments	Strongest direct network effect lever	Supports supplier proofing in cross-border trade
C10 Financing	Faster financing readiness checkpoints	Keeps teams using Stoquify for decision evidence	Referrals to partner financing channels	Lender readiness/partner transaction fee upside	Consent and evidence controls reduce bias/abuse risk	Reduced manual underwriting support burden	Creates auditable financing evidence pack dataset	Integrates financing posture into existing workflows	Partner network moat with documented consent pathways	Enhances credit evidence maturity for OHADA-compliant lending docs
C11 Pipeline-to-Cash	Faster quote-to-cash progression	Keeps revenue process embedded in operations	Increases referrals via better client response	Revenue acceleration and cash visibility features	Evidence-based commitments reduce misunderstandings	Improves conversion without overstaffing sales ops	Adds stage and conversion history for forecasting	Links sales, collections and finance workflows	Expands account-level growth loops	Supports revenue recognition and reporting consistency
C12 Project Margin	Faster project risk and margin checks	Increases repeat use for project-heavy clients	Shared estimates and supplier alignment drives referral trust	Verticalized pricing/estimation products	Evidence and assumptions make disputes easier to resolve	Prevents margin shocks and rework costs	Enriches project-time-cost dataset for next period	Creates consistent closeout and change-control workflow	Deepens subcontractor and partner reliance on shared process	Aligns project profitability evidence with OHADA bookkeeping rigor

Programme north star and guardrails

North star: monthly organizations completing at least three evidence-backed business outcomes across at least two Stoquify workflows, with no material trust-policy violation.

Dimension	Leading indicator	Lagging indicator	Guardrail
Daily use	Outcome days per active user/week	8-week retained active organizations	Do not count passive notifications as value
Retention	Multi-workflow activation; second user	Gross/net revenue retention	Exportability and revocation remain available
Referral	Invitations and acceptance	Referred activated/paying organizations	Consent and complaint rate
Revenue	Paid feature activation	Expansion ARR and gross margin	Cost per verified outcome
Trust	Evidence coverage; approval comprehension	Critical incidents and disputed outputs	Zero tolerated cross-tenant leakage
Operations	Automated detection and successful recovery	Support cost and MTTR	No hidden retry or duplicate action

12. Staffing and operating model

12.1 Recommended team

Function	Recommended FTE	Minimum FTE	Accountability
Programme/product leadership	2–3	1–2	Portfolio, discovery, value gates
Platform/runtime/backend	4–6	3	Orchestration, APIs, cases, service contracts
Security/privacy/identity	2–3	1	Threat models, authorization, consent, incident readiness
Data/ML/AI platform	3–4	2	Evidence, models, evaluation, forecasting/signals
Frontend/mobile/design	3–4	2	Queues, collaboration, verification and accessibility
QA/assurance/SRE	3–4	2	Automation, release gates, observability and operations
Domain experts	2–3 shared	1 shared	Accounting, collections, OHADA/country rules
Partnerships/growth/success	2–3	1	Providers, advisors, pilots, packaging and adoption
Total	21–30	13–16	Some roles may be fractional early

The 16–22-person delivery assumption used in the schedule reflects ramping and shared/fractional roles. The minimum team must reduce concurrency; it must not remove security, domain review or independent assurance.

12.2 RACI matrix

Workstream	Responsible (R)	Accountable (A)	Consulted (C)	Informed (I)
Program baseline & governance	Programme Lead	CTO/CPO	Assurance Authority, Security Lead	All leads
Runtime foundations (F-A to F-G)	Platform Lead	CTO	SRE, Security, Product	Programme Lead
Shared evidence, consent, trust graph	Security/Privacy Lead	Legal Counsel	Product, SRE, Data Ops	Assurance Authority
C01 Collections + Advisor + Document intake	Product Lead (Collections) + Domain Product Manager	Programme Lead	Compliance, Finance, Frontend	Customer Success
C04 Forecast + C07 Profitability	Data Finance Lead	Programme Lead	Controller/Accounting, Product	Sales Lead
C05 Communications + C06 Controls + C10 readiness	Security & Product Lead	Compliance Officer	Communications, Legal, Partner team	Product Lead
Supplier + Project + Pipeline-to-Cash modules	Program Manager	Programme Lead	Sales, Supplier Partnerships, Growth	Finance Lead
Evaluation, release, canary, rollbacks	QA/Assurance Lead	CTO/SRE Lead	Product, Security, Programme Lead	All domain leads

12.3 Operating cadence

- Weekly capability delivery and evidence review.
- Fortnightly architecture/security council for contract and autonomy changes.
- Monthly value, cost and trust review with pilot owners.
- Quarterly continue/pivot/stop portfolio decision.
- Incident review within 48 hours for material Copilot, privacy or external-action events.
- Country/provider release is separately versioned; approval in one jurisdiction does not authorize another.

13. Commercial and financial plan

Packaging hypothesis

- **Core:** meaningful document allowance, connector trust, read-only insights and basic advisor invitation.
- **Cash Operations:** Collections, cash forecast and communications usage.
- **Finance Professional:** Advisor cockpit, client bundles, close/evidence collaboration.
- **Controls and Performance:** Continuous controls, profitability and advanced assurance.
- **Network/Partner:** Supplier collaboration and financing-readiness packages, with transparent partner terms.
- **Vertical:** Pipeline-to-Cash and Project Margin packages tailored to chosen industries.

Do not finalize pricing before measuring willingness to pay, cost per verified outcome and support burden. Avoid charging external recipients merely to respond to a request; friction would weaken the referral loop.

Investment checkpoints

Checkpoint	Release next funding when...	Stop/piv ot when...
After Phase 0 owners, prerequisites, baseline and architecture decisions are credible	no accountable release authority or pilot cohort exists	
After Phase 1 base runtime and critical evaluation gate are green		tenant isolation or side-effect controls remain structurally unsafe
After Phase 2 top-three pilots show task value, acceptable cost and trust		accuracy/complaints/support costs exceed agreed budgets
After Phase 3 forecasts beat baseline and control signals are useful		users over-trust outputs or false positives remain harmful
After Phase 4 external-network activation and partner economics are proven		legal/consent/fairness risks cannot be bounded
After Phase 5 one revenue/vertical wedge has repeatable economics		generic CRM or broad vertical scope dilutes core advantage

14. Risk register

ID	Risk	Likelihood/impact	Prevention	Detection/response	Owner
R01	Cross-tenant/client data leakage	M/Critical	Trusted context, resource grants, least privilege	Continuous denial tests; kill switch; incident response	Security
R02	Unsafe or duplicate external message	M/High	Approval, idempotency, consent, templates	Delivery reconciliation, complaints, immediate suspension	Communications
R03	Model fabricates financial/compliance truth	M/High	Deterministic services, structured outputs, evidence	Unsupported-claim tests and user dispute flow	AI platform
R04	Stale/broken connector contaminates advice	H/High	C08 trust contract and freshness gates	Trust banner, blocked output, incident queue	Integration
R05	Document malware/PII breach	M/Critical	Quarantine, scan, encryption, retention	Security monitoring and breach runbook	Security
R06	False fraud accusation	M/High	Signal wording, deterministic basis, restricted review	Precision monitoring, correction and appeal	Controls
R07	Forecast false precision	H/Medium	Assumption ledger/confidence, backtests	Error-by-horizon dashboard and degradation	Finance data
R08	External identity/relationship abuse	M/Critical	Purpose/resource/time scopes, fresh auth	Access anomaly and rapid revocation	Identity
R09	Financing privacy/fairness violation	M/Critical	Consent, minimization, partner/legal review	Fairness/complaint audits, suspend partner	Privacy/legal
R10	Provider/country coverage fragmentation	H/Medium	Adapter contracts and staged countries	SLA/incident metrics and fallback	Partnerships
R11	Cost/latency runaway	M/High	Routing budgets, caching, deterministic fallback	Per-outcome FinOps alerts and throttles	AI/SRE
R12	Scope expansion into generic CRM/ERP	H/High	Narrow outcome charters and gates	Quarterly stop/pivot review	Product
R13	Evaluation theater	M/Critical	Runnable fixtures and immutable artifacts	Independent gate and reproducibility audit	Assurance
R14	External prerequisites delay critical path	H/High	named owners/dates/escalation	weekly blocker aging and alternate design work	Programme lead
R15	Low advisor/supplier adoption	M/High	lightweight portal, free response, co-design	funnel/task telemetry and interviews	Growth
R16	Data residency/retention mismatch	M/High	country data map and configurable retention	compliance scan and deletion proof	Privacy

15. First 30, 60 and 90 days

Days 0–30 — make the programme governable

1. Appoint executive sponsor, programme lead, platform owner, security owner, assurance authority and product owners for C01/C02/C03/C08.
2. Create the authoritative status register and rerun the current runtime gates against an immutable candidate.
3. Turn all external-input blockers into named tasks with owner, evidence requested, due date and escalation path.
4. Freeze autonomy policy at L0–L2; publish the prohibited-action catalog.
5. Complete discovery with 8–12 businesses, 4–6 accountants/advisors, and AR/document operators in the first target countries.
6. Baseline days-to-pay, document cycle time, evidence completeness, advisor invitation behavior, connector incidents and current support cost.
7. Approve ADRs for case lifecycle, external-party identity, consent, evidence envelope and provider abstraction.

Days 31–60 — build the proof foundations

1. Implement the executable evaluation harness and highest-risk 150 cases.
2. Close or evidence production target, CI artifact identity, workload identity, secret handling and operational collectors.
3. Build common connector-health and evidence-envelope contracts with two existing provider adapters.
4. Prototype the verification inbox, get-paid queue and advisor request queue with clickable evidence/approval states.
5. Implement synthetic document quarantine/scan/extract/verify flow and a deterministic collections priority fixture.
6. Finalize pilot tenant selection, data-processing terms, on-call roster, stop conditions and success thresholds.

Days 61–90 — demonstrate controlled end-to-end slices

1. Execute all 999 baseline cases or publish an explicit blocked exception; no untested case may be reported as passed.
2. Run base Command Agent Phase 2B entry and controlled-pilot rehearsal, including suspension and rollback.
3. Demonstrate receipt/photo → quarantine → extract → human verify → AP draft with complete provenance.
4. Demonstrate overdue invoice → evidence-linked priority → approved unsent message preview → promise case using synthetic data.
5. Demonstrate advisor invitation → resource-scoped request → evidence response → revocation with cross-client denial tests.
6. Present the steering committee with measured accuracy, task completion, cost, security and readiness evidence; authorize or reject live canaries separately.

16. Final recommendations by capability

Capability	Decision	Start now	Delay until	Explicit rejection boundary
C01 Collections	START design/build	AR contracts, rules, drafts	outbound send until consent/runtime gates	Autonomous legal threats, discounts or restructuring
C02 Document Intake	START	secure pipeline, two document classes	silent posting and sensitive classes	Unverified extraction posted as truth
C03 Advisor	START	identity design, requests, read-only cockpit	broader firm powers until denial suite	Cross-client discovery or ambiguous certification
C04 Forecast	PREPARE	source contracts/backtests	user pilot until freshness foundation	Model-generated booked facts
C05 Communications	PREPARE	consent/templates/provider design	sending until C01/C03 and provider gate	Unapproved mass or consequential messaging
C06 Controls	PREPARE	narrow rules/confidential cases and deterministic signals	statistical expansion until outcome labels	Automated accusation, discipline or payment block
C07 Profitability	PREPARE	allocation engine and data-quality study	price drafts until costs reconcile	Automatic price publication
C08 Connector Health	START as foundation	common health model and two adapters	auto-remediation of credentials	Credential exposure to models
C09 Supplier	DELAY product pilot	research and authorization reuse	C03 external-access proof	Unverified bank-detail change
C10 Financing	DELAY	legal/partner discovery and schemas	mature consent/evidence and regulated partner	Lending decision, hidden score or data sale
C11 Pipeline-to-Cash	DELAY	narrow discovery	C01+C05+C07 proven	Generic CRM rebuild or unapproved commitments
C12 Project Margin	DELAY	select one vertical	C07 and project/time data ready	Cross-industry generic build

17. Verification record and evidence standard

This roadmap was reconciled against:

- the authoritative Top 12 research report;
- package.json runtime/release scripts;
- current services/agents, actions/agents and Copilot/Command UI surfaces;

- `prisma/schema.prisma` agent, invitation, provider-event and evidence-related models;
- the July phased execution completion audit, blocker-elimination/Phase 3 entry report, operational release evidence, external-input readiness and Phase 3 decision;
- the 14 June architecture graph report.

Required verification commands for programme execution include:

```
npm run prisma:validate
npm run agent:tool-registry:gate
npm run agent:prohibited-action:gate
npm run agent:release-control:gate
npm run ai:copilot:guardrails:gate
npm run agent:phase2a:gate
npm run agent:phase2b:entry:gate
npm run agent:phase3:entry:gate
```

Each command result must be bound to source commit, immutable artifact, environment class and timestamp. A local pass is development evidence; it is not a substitute for production-environment prerequisites, independent approval or a controlled pilot.

18. Latest verification run log

Latest executable verification (this run, 2 August 2026)

Command	Result	Exit code	Duration
python docs/copilot/stoquify-agent-skill-definition-suite/scripts/validate_suite.py	PASS (agents=9, skills=28, capabilities=37, cases=999)	0	3.2s
npm test -- --runInBand services/agents/portfolio/__tests__/connector-inventory-read-model.service.test.ts			
services/agents/portfolio/__tests__/connector-health.service.test.ts	PASS (3 suites, 15 tests)	0	26.1s
services/agents/portfolio/__tests__/evidence-trust.contracts.test.ts			
npm run typecheck	PASS	0	34.6s
npm run prisma:validate	PASS	0	11.1s
npm run agent:tool-registry:gate	PASS	0	3.4s
npm run agent:prohibited-action:gate	FAIL (services/agents/portfolio/connector-inventory-read-model.service.ts Prisma DB import outside governance boundary)	1	3.9s
npm run agent:release-control:gate	PASS	0	4.3s
npm run ai:copilot:guardrails:gate	PASS	0	3.4s
npm run agent:phase2a:gate	PASS	0	3.4s
npm run agent:phase2b:entry:gate	BLOCKED (2/23 passed, 21 blockers)	1	3.6s
npm run agent:phase3:entry:gate	BLOCKED (0/34 passed, 34 blockers)	1	3.5s
npm run agent:external-inputs:report	EXTERNAL_INPUTS_REQUIRED (1/13, 102 blockers)	0	3.6s

All command results are commit-stamped and used as part of Phase 0 evidence; FAIL/BLOCKED items are intentionally treated as hard guards, not defects.

19. Approval requested

Leadership should approve:

1. Phase 0 and Phase 1 funding and named accountability.
2. The top-three value loop plus Connector Health as the first delivery portfolio.
3. The L0–L2 autonomy ceiling and prohibited-action boundaries.
4. A gate-based 24–30 month programme envelope, reviewed quarterly.
5. External prerequisite acquisition—production target, identities, secrets, approvers, statutory reviewers, provider credentials and pilot organizations—as first-class programme work.

The immediate next move is not to write twelve prompts. It is to close the runtime qualification gap while producing the shared contracts and three end-to-end evidence-backed pilot slices that prove Stoquify can safely turn trusted business data into repeated outcomes, collaboration and referrals.